

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/277288211>

# GoF Structural Patterns: A Formal Specification

Article · January 2000

---

CITATIONS

7

---

READS

143

2 authors, including:



[Andrés Pablo Flores](#)

National University of Comahue

37 PUBLICATIONS 222 CITATIONS

SEE PROFILE

Some of the authors of this publication are also working on these related projects:



Domain-Oriented Software Reuse [View project](#)



The United Nations  
University

**UNU/IIST**

International Institute for  
Software Technology

---

# GoF Structural Patterns: A Formal Specification

---

Andres Pablo Flores and Richard Moore

August 2000

## UNU/IIST and UNU/IIST Reports

UNU/IIST (United Nations University International Institute for Software Technology) is a Research and Training Centre of the United Nations University (UNU). It is based in Macau, and was founded in 1991. It started operations in July 1992. UNU/IIST is jointly funded by the Governor of Macau and the governments of the People's Republic of China and Portugal through a contribution to the UNU Endowment Fund. As well as providing two-thirds of the endowment fund, the Macau authorities also supply UNU/IIST with its office premises and furniture and subsidise fellow accommodation.

The mission of UNU/IIST is to assist developing countries in the application and development of software technology.

UNU/IIST contributes through its programmatic activities:

1. Advanced development projects, in which software techniques supported by tools are applied,
2. Research projects, in which new techniques for software development are investigated,
3. Curriculum development projects, in which courses of software technology for universities in developing countries are developed,
4. University development projects, which complement the curriculum development projects by aiming to strengthen all aspects of computer science teaching in universities in developing countries,
5. Courses, which typically teach advanced software development techniques,
6. Events, in which conferences and workshops are organised or supported by UNU/IIST, and
7. Dissemination, in which UNU/IIST regularly distributes to developing countries information on international progress of software technology.

Fellows, who are young scientists and engineers from developing countries, are invited to actively participate in all these projects. By doing the projects they are trained.

At present, the technical focus of UNU/IIST is on formal methods for software development. UNU/IIST is an internationally recognised center in the area of formal methods. However, no software technique is universally applicable. We are prepared to choose complementary techniques for our projects, if necessary.

UNU/IIST produces a report series. Reports are either Research R, Technical T, Compendia C or Administrative A. They are records of UNU/IIST activities and research and development achievements. Many of the reports are also published in conference proceedings and journals.

Please write to UNU/IIST at P.O. Box 3058, Macau or visit UNU/IIST home page: <http://www.iist.unu.edu>, if you would like to know more about UNU/IIST and its report series.

Zhou Chaochen, Director — 01.8.1997 – 31.7.2001



The United Nations  
University

**UNU/IIST**

International Institute for  
Software Technology

P.O. Box 3058  
Macau

---

# GoF Structural Patterns: A Formal Specification

---

Andres Pablo Flores and Richard Moore

## Abstract

GoF structural patterns represent standard techniques in object-oriented design of composing classes and objects to form larger structures. The GoF catalogue describes, using a standard but informal notation, seven such patterns, each of which captures different structural aspects. In this paper, we present an analysis of the essential components and properties of these patterns, and we specify these properties formally using a formal model of a general object-oriented design which was developed in earlier work as the basis for the specification. We also give an example of how to use these specifications in order to check whether a subset of a particular design matches a particular structural pattern.

Andres Pablo Flores is a Fellow of UNU/IIST (November 1999 to August 2000), on leave from Comahue University, Neuquen, Argentina, where he is an Assistant Teacher. His research interests are focused on the software engineering disciplines, mainly on those related with software analysis and design. He is currently working on the combination of formal and informal methods and its application in a real domain.

Richard Moore is a Research Fellow on the staff of UNU/IIST, a position he took up on October 1st 1995. He has an M.A. in mathematics from the University of Cambridge and a Ph.D. in physics from the University of Manchester. He has been engaged in computing science research in the field of formal methods since 1985, a large part of which was carried out in the formal methods group at Manchester University. He has written several papers on formal methods and is co-author of two books on formal methods – mural: a Formal Development Support System; and Proof in VDM: A Practitioner's Guide. He has also worked for the Defence Research Agency in Malvern, UK, on various formal methods projects, both as a consultant and as a full-time member of staff.

## Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>The Adapter Pattern</b>	<b>3</b>
2.1	Properties of the Adapter Pattern . . . . .	4
2.2	Formalising the Adapter Pattern . . . . .	6
2.2.1	Formal specification of the Class Adapter pattern . . . . .	7
2.2.2	Formal specification of the Object Adapter pattern . . . . .	12
<b>3</b>	<b>The Bridge Pattern</b>	<b>16</b>
3.1	Properties of the Bridge Pattern . . . . .	16
3.2	Formalising the Bridge Pattern . . . . .	17
<b>4</b>	<b>The Composite Pattern</b>	<b>22</b>
4.1	Properties of the Composite Pattern . . . . .	23
4.2	Formalising the Composite Pattern . . . . .	24
<b>5</b>	<b>The Decorator Pattern</b>	<b>32</b>
5.1	Properties of the Decorator Pattern . . . . .	32
5.2	Formalising the Decorator Pattern . . . . .	34
<b>6</b>	<b>The Facade Pattern</b>	<b>40</b>
6.1	Properties of the Facade Pattern . . . . .	41
6.2	Formalising the Facade Pattern . . . . .	42
<b>7</b>	<b>The Flyweight Pattern</b>	<b>48</b>
7.1	Properties of the Flyweight Pattern . . . . .	49
7.2	Formalising the Flyweight Pattern . . . . .	51
<b>8</b>	<b>The Proxy Pattern</b>	<b>58</b>
8.1	Properties of the Proxy Pattern . . . . .	58
8.2	Formalising the Proxy Pattern . . . . .	60
8.3	Virtual Proxy . . . . .	65
<b>9</b>	<b>Application of the Decorator Pattern</b>	<b>67</b>
<b>10</b>	<b>Conclusions</b>	<b>72</b>
<b>A</b>	<b>Specification of the Decorated Text Editor</b>	<b>73</b>



# 1 Introduction

Designing is a difficult task. The efficiency of the design process needs to be improved in order to avoid wasting time and effort which only translate to more costly software development.

One important step in this direction is to reuse design work that has already been done, and which is therefore proven to work well and can thus serve as certification of a new design. This requires abstracting out pieces of design that have been found to recur in previous designs. This is still a very hard task, however, even for expert designers.

One approach to reuse in which there is widespread interest is design patterns. Each design pattern describes a generic solution to some problem which is commonly encountered in different contexts, and the solution to a particular problem can be obtained by customising the generic pattern design [8, 16]. Patterns thus summarise the experience of designers working on similar problems in different contexts and represent proven solutions for solving these problems.

The solution proposed by a pattern involves a sort of structure which properly balances the numerous competing concerns or “forces” which are present in a certain context [1]. Design patterns convey regularities, plans, aspects, or abstractions of programs rather than concrete instances. Applying a design pattern allows some aspects in a system structure to vary independently of others, emphasising flexibility in the whole system, because the structure of the pattern results in a high degree of adaptability [15]. In addition, patterns are described in a general and abstract way which helps designers recognise an architecture matching the pattern when working on a particular concrete application. Thus, most design patterns have an unlimited number of implementations, usually in various programming languages and numerous application domains [5]. Their use can therefore lead to a more rapid understanding of a particular design problem and hence a more rapid completion of the design, which means saving time and effort in the whole development [4].

Design patterns became popular with the publication of the GoF catalogue<sup>1</sup> [7], which introduced a body of literature for design problems in software development, similar to the common vocabularies which are fundamental in any science or engineering discipline for expressing and relating its concepts [1]. The patterns thus help to create a shared language for communicating insight and experience about particular design problems and their possible solutions.

GoF design patterns are described by means of natural language narrative and a graphical notation which is based on an extension of OMT (Object Modelling Technique [13]). Although these are good tools for representing the essence of each pattern intuitively, it is not sufficiently precise to allow a designer to demonstrate conclusively that a particular problem matches a particular pattern or that a proposed solution is consistent with a particular pattern. The notation also makes it difficult to be certain that the patterns used are meaningful and contain no inconsistencies. It is therefore extremely difficult to give any meaningful certification of the correctness of software developed using patterns.

---

<sup>1</sup>GoF means Gang of Four, as the authors of this catalogue are commonly referred to.



Providing a more precise description of patterns can help designers know more clearly not only when and how a particular pattern can be applied but also that it has been applied correctly. It can also help to improve understanding of the patterns and to avoid inconsistencies, ambiguities and incompleteness which are inherent in the graphical/textual notation.

To this end, a formal model of a generic object-oriented design has been developed [6] and specified using the RAISE specification language RSL [11]. This formalises the various components which are found in the extended OMT notation used in the GoF catalogue and also separates the design from the patterns. Various common properties of the GoF design patterns are then specified in this model as generic RSL functions, and these, appropriately instantiated and combined, can be used to formalise the properties of the patterns. The verification that a design matches a pattern is then done by first relating or binding the names of the entities (classes, methods, state variables, and parameters) in (a subset of) the design with the names of corresponding entities appearing in the particular pattern, then checking that all the properties of the pattern are satisfied by the corresponding entities in the design.

Figure 1 illustrates how the *renaming map* binds a (subset of a) design to a particular pattern – the bindings are denoted by the dotted lines. As can be seen in this figure, there can be some flexibility in the way these bindings are formed. For example, a class hierarchy in the pattern can correspond to a more complicated hierarchy in the design where there are intermediate classes which play no role in the pattern, and operations can be defined or implemented in a superclass of the class corresponding to the one in which they appear in the pattern.

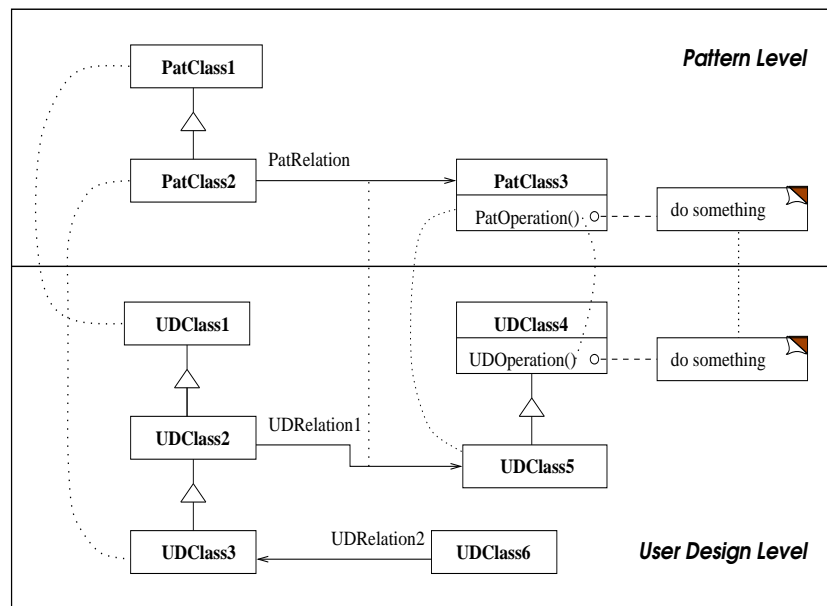


Figure 1: Binding of Design and Pattern Levels

GoF design patterns are classified in [7] according to two criteria: *purpose* and *scope*. A pattern's purpose reflects what the pattern does, and may be either *creational*, *structural*, or *behavioural*, while the scope indicates whether the pattern applies primarily to classes or to objects, and thus

may be either *class* or *object*.

Nine of the eleven behavioural patterns have already been specified as described above [12], and parallel work is addressing the creational patterns [2]. In this report we apply the same techniques to the GoF structural patterns.

Structural patterns are concerned with how classes and objects are composed to form larger structures. In general, structural class patterns use inheritance to compose interfaces or implementations, while structural object patterns describe ways to compose objects so as to realize new functionality. Object composition gives some added flexibility because the composition can be changed at run-time, which is impossible with static class composition, so the scope of most structural patterns is in fact object.

Patterns in general encapsulate aspects of a design that can vary, with individual patterns relating to the variation of different aspects. A designer can use these aspects to help determine which is the most appropriate pattern for a particular situation. Table 1 lists the design aspect(s) addressed by the various structural design patterns.

Design pattern	Aspect(s) that can vary
Adapter	interface to an object
Bridge	implementation of an object
Composite	structure and composition of an object
Decorator	responsibilities of an object without subclassing
Facade	interface to a subsystem
Flyweight	storage costs of objects
Proxy	how an object is accessed; its location

Table 1: Design Aspects of Structural Patterns

In Sections 2 to 8 we present discussions and specifications of the properties of each of the GoF structural patterns, then in Section 9 we illustrate, using an example of a design based on the Decorator pattern, how our model can be used to check whether or not a given design matches a given pattern. We end with a summary of our work and an indication of possible future work.

## 2 The Adapter Pattern

The Adapter pattern allows the reuse of existing tools even though their interface may not coincide with that of some of the classes that need to use them. There are two ways in which a relationship between the different interfaces can be established, for each of which there is a corresponding form of the pattern.

The first uses inheritance relations, specifically multiple inheritance, so the relationship is established statically. This form is classified with the scope **class** in the GoF catalogue [7]. The second form, which has scope **object**, uses static and dynamic relations instead to establish the relationship. We first introduce and discuss the properties of the pattern as described in [7], then we describe our formalisation of the two forms of the pattern using our generic formal model [6].

## 2.1 Properties of the Adapter Pattern

The general properties of the Adapter pattern are defined in [7] as follows:

### Intent

Convert the interface of a class into another interface clients expect. Adapter lets classes work together that couldn't otherwise because of incompatible interfaces.

### Applicability

Use the Adapter pattern when:

- you want to use an existing class, and its interface does not match the one you need.
- you want to create a reusable class that cooperates with unrelated or unforeseen classes, that is, classes that don't necessarily have compatible interfaces.
- (*object adapter only*) you need to use several existing subclasses, but it's impractical to adapt their interface by subclassing every one. An object adapter can adapt the interface of its parent class.

### Structure

A class adapter uses multiple inheritance to adapt one interface to another, while an object adapter uses object composition. The structures of the two forms of the pattern are shown in Figures 2 and 3 respectively.

### Participants

- Target
  - defines the domain-specific interface that Client uses.
- Client
  - collaborates with objects conforming to the Target interface.
- Adaptee
  - defines an existing interface that needs adapting.
- Adapter
  - adapts the interface of Adaptee to the Target interface.

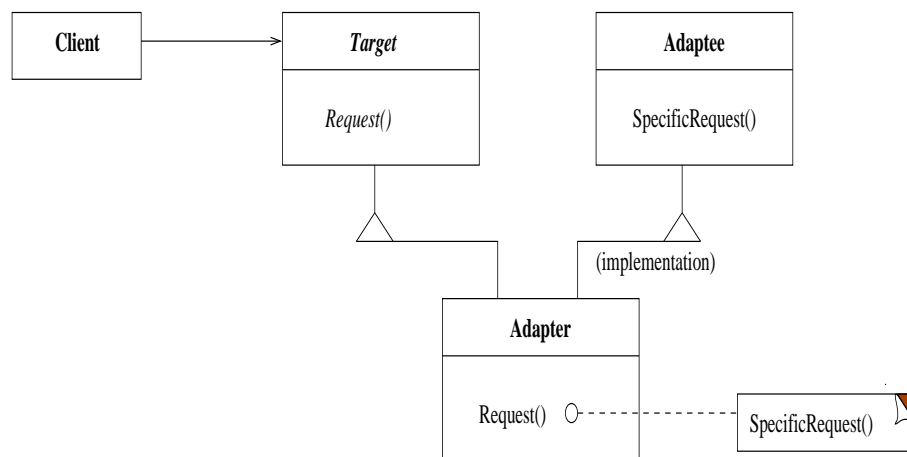


Figure 2: Class Adapter Pattern Structure

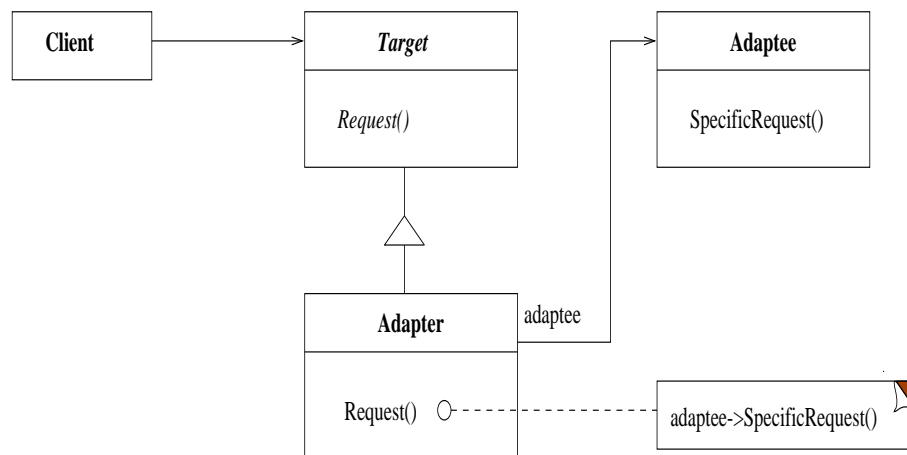


Figure 3: Object Adapter Pattern Structure

## Collaborations

- Clients call operations on an Adapter instance. In turn, the adapter calls Adaptee operations that carry out the request.

## 2.2 Formalising the Adapter Pattern

As seen above, the two different forms of the pattern have different structures, and there are important differences between their properties. We therefore analyse and specify each of the forms separately. However, some features are common to both forms and these are only explained once.

One consideration that is relevant to both forms of the pattern is that the structure shows the **Adapter** class to be the only subclass of the **Target** class. While this represents one possible implementation of the Adapter pattern, corresponding to a design in which the **Target** class effectively unifies the interfaces of many different and disparate classes using a number of different **Adapter** classes, there is another possible implementation in which the **Target** class itself defines the interface for a hierarchy of **ConcreteTarget** subclasses and the **Adapter** class is used to make the interface of the **Adaptee** conform to this existing common interface.

We make this second possibility explicit in our treatment of both forms of the Adapter pattern by introducing the new role **ConcreteTarget** into the structure, though we make no assumptions about this except that it implements the **Request** interface of the abstract **Target** class. We similarly introduce the role **SubclassAdaptee** to represent possible subclasses of the **Adaptee** class which do not play the **Adapter** role, though in this case the new role is only introduced in the Class Adapter because in the Object Adapter the **Adaptee** class does not belong to a hierarchy so its subclasses are irrelevant. This leads us to the modified pattern structures for the Class and Object Adapter patterns shown in Figures 4 and 5 respectively.

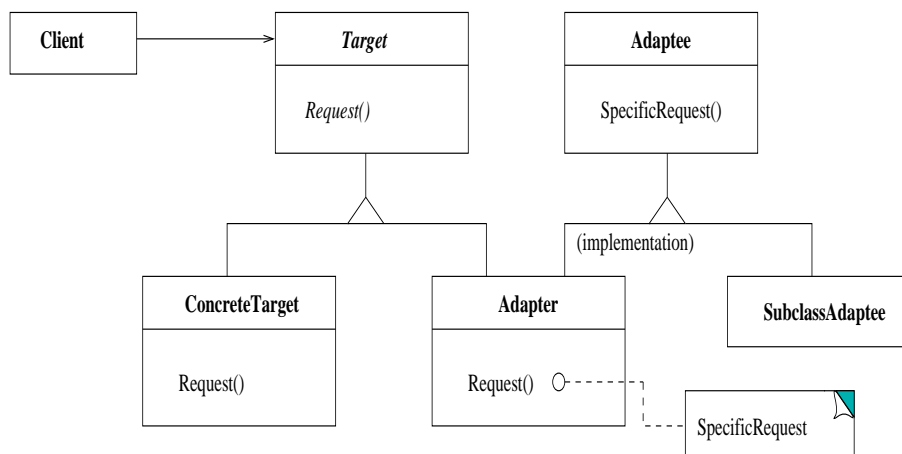


Figure 4: The Modified Class Adapter Pattern Structure

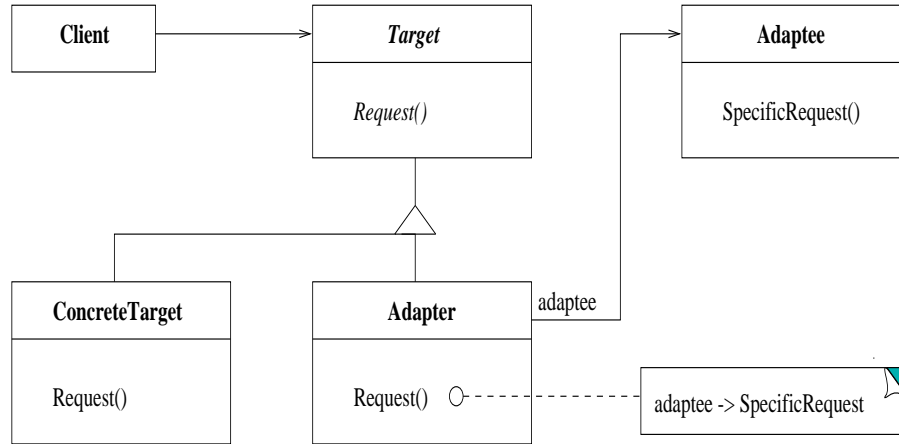


Figure 5: The Modified Object Adapter Pattern Structure

Based on these modified structures, we now proceed to give the formal specifications of the two forms of the pattern.

### 2.2.1 Formal specification of the Class Adapter pattern

We begin by defining RSL constants representing the names of all of the different entities (classes and methods) appearing in the modified pattern structure shown in Figure 4:

**value**

Target, ConcreteTarget, Adapter, Adaptee, SubclassAdaptee, Client : G.Class\_Name,

Request, SpecificRequest : G.Method\_Name

We now define the properties of the pattern by defining the properties that each of the pattern entities defined above must satisfy. This is done using the functions representing the generic properties of patterns which are defined in [6].

Consider first the hierarchy of classes consisting of the abstract class **Target** and its concrete subclasses **ConcreteTarget** and **Adapter**. Since the **Target** class defines the interface the **Client** class uses, we can consider it to be unique without any loss of generality – a second **Target** class would correspond to a different interface and we can consider this situation to be two distinct instances of the pattern. In addition, the classes at the leaves of the hierarchy must play either the **ConcreteTarget** or the **Adapter** role (but not both because the **ConcreteTarget** role was introduced to represent classes in the hierarchy other than the **Adapter** class) and no class in the hierarchy can play the **Client** role (because the **Target** class represents the interface used by the **Client**) or either the **Adaptee** or the **SubclassAdaptee** role (because that would mean the

adaptee classes inherit the interface of the **Target** class, in which case the Adapter pattern would not be needed).

We also take account of the design heuristic which suggests that when related classes have properties in common these properties should be factored out into a superclass, the so-called process of classification or generalisation [9, 13]. Thus, in a design corresponding to the Adapter pattern we may have a class playing the **ConcreteTarget** role which has subclasses which also play the same role, the subclasses representing specialisations of the concrete target classes defined by the superclass. We therefore generalise the inheritance relation to allow intermediate classes in the hierarchy, that is additional classes between the root and the leaves, which can play the same roles as the leaves, that is **ConcreteTarget** or **Adapter**. However, we impose a constraint that requires that if an intermediate class plays such a role it must play the same role as all its leaf subclasses. This means that one **ConcreteTarget** class can be a specialisation of another but cannot be a specialisation of an **Adapter** class.

All these properties together are specified using the generic function *hierarchy* from [6], instantiated with the appropriate roles:

```

value
  one_target_in_hierarchy : Wf_Design_Renaming → Bool
  one_target_in_hierarchy(dr) ≡
    hierarchy
      (Target, {ConcreteTarget, Adapter}, {Client, Adaptee, SubclassAdaptee}, dr)

```

Similar considerations apply to the hierarchy of classes beginning from the **Adaptee** class – again we can consider the class to be unique without loss of generality, interpreting a design in which there are two classes playing the **Adaptee** role as two distinct instances of the pattern (possibly with the same **Client** or **Target** classes), and the properties of the subclasses **Adapter** and **SubclassAdaptee** are entirely analogous to those of **Adapter** and **ConcreteTarget** described above. We therefore use the same function *hierarchy* to formalise these properties, though instantiating it with different roles of course:

```

value
  one_adaptee_in_hierarchy : Wf_Design_Renaming → Bool
  one_adaptee_in_hierarchy(dr) ≡
    hierarchy
      (Adaptee, {Adapter, SubclassAdaptee}, {Client, Target, ConcreteTarget}, dr)

```

We have already specified above that both the **Target** and **Adaptee** classes are unique, that is there should only be one class playing each of those roles in the design. As far as the **Adapter** role is concerned, however, it is possible that the **Client** needs to adapt the **Adaptee**'s behaviour in several different ways, that is it requires several different implementations of the operations

included in the *Adaptee*'s interface. In this case more than one *Adapter* class may be required in the design, each one adapting the behaviour of the *Adaptee* class in a different way. Thus, we do not impose the restriction that there should be only one *Adapter* class, though we do insist that there must be at least one class in the design playing this role otherwise the design does not implement the Adapter pattern. The generic function *exists\_role* from [6], appropriately instantiated with the role *Adapter*, is used to specify this property:

```
value
  exists_adapter : Wf_Design_Renaming → Bool
  exists_adapter(dr) ≡ exists_role(Adapter, dr)
```

Regarding the *ConcreteTarget* and *SubclassAdaptee* classes, these were introduced to take account of the fact that both the *Target* and *Adaptee* classes might (though do not necessarily) have subclasses other than the *Adapter* class. Classes playing these roles therefore may or may not exist in the design so we leave the existence of these classes unspecified.

One thing we can say about these classes, however, is that if they do exist they must be concrete subclasses of the *Target* and *Adaptee* classes respectively. This is specified using the generic function *is\_concrete*. This property additionally ensures that the *Request* methods in the *ConcreteTarget* classes must be concrete – according to our model of a general object-oriented design [6], a concrete class cannot contain abstract methods so the abstract *Request* method inherited from the *Target* class must be overridden by a concrete version.

```
value
  CT_is_concrete_target : Wf_Design_Renaming → Bool
  CT_is_concrete_target(dr) ≡ is_concrete(Target, ConcreteTarget, dr),

  subclassAdaptee_is_concrete : Wf_Design_Renaming → Bool
  subclassAdaptee_is_concrete(dr) ≡ is_concrete(Adaptee, SubclassAdaptee, dr)
```

The same generic function is also used to specify that the *Adapter* role is concrete and is a subclass both of the *Target* class and of the *Adaptee* class:

```
value
  Ar_has_two_parents : Wf_Design_Renaming → Bool
  Ar_has_two_parents(dr) ≡
    is_concrete(Target, Adapter, dr) ∧
    is_concrete(Adaptee, Adapter, dr)
```

The *Adaptee* class should also be concrete. However, although it has a concrete interface this does not automatically mean that it is a concrete class because an abstract class can also have



a concrete interface. We therefore need to express explicitly the fact that the `Adaptee` class is concrete. This is done using the function `is_concrete_class` from [6]:

```

value
  is_concrete_adaptee : Wf_Design_Renaming → Bool
  is_concrete_adaptee(dr) ≡ is_concrete_class(Adaptee, dr)

```

The `Target` class should be abstract. However, its interface includes at least one `Request` operation, which is also abstract, and this automatically means that the class must be abstract – one of the properties of a general object-oriented design which is incorporated into our general model in [6] is that a class containing abstract methods must itself be abstract. We therefore do not need to specify explicitly that the `Target` class is abstract. Rather we simply specify that it contains an abstract `Request` method. This is done using the function `has_def_method`.

Of course the `Target` class can contain more than one `Request` method – these represent the interface which is used by the `Client` and which is to be adapted by the `Adapter` and this interface can clearly comprise more than one method. We therefore also specify that all methods in this interface should be abstract, which is checked using the function `has_all_def_method`. We do not specify any other properties of these methods, however, because we have no a priori knowledge about their implementation – different methods may have different bodies and may or may not have results and parameters.

The full specification of the properties of the `Request` methods in the `Target` class is therefore as follows:

```

value
  T_has_defined_request : Wf_Design_Renaming → Bool
  T_has_defined_request(dr) ≡
    has_def_method(Target, Request, dr) ∧
    has_all_def_method(Target, Request, dr)

```

The `SpecificRequest` methods in the `Adaptee` class are treated similarly except that these methods are implemented instead of abstract. We therefore use the functions `has_impl_method` and `has_all_impl_method`, which are analogous to `has_def_method` and `has_all_def_method` except that they check respectively that at least one implemented method playing a particular role exists and that all methods playing the given role are implemented. Again, we do not specify any other properties of these methods because there may be a number of different methods playing the same role but having different bodies, different results and different parameters.

```

value
  Ae_has_specific_request : Wf_Design_Renaming → Bool

```

```

Ae_has_specific_request(dr) ≡
  has_impl_method(Adaptee, SpecificRequest, dr) ∧
  has_all_impl_method(Adaptee, SpecificRequest, dr)

```

The existence and concreteness of the `Request` methods in the `Adapter` class follows from the concreteness of the class just as in the case of the `ConcreteTarget` class, but in this case the implementation of the methods is also fixed by the pattern – each `Request` method simply performs an invocation to one of the `SpecificRequest` methods which the `Adapter` class inherits from the `Adaptee` class. In the general model [6] this corresponds to a *self-invocation* and uses the reserved variable *self* to specify the invocation of some method in the interface of the same class. This is described using the generic function *self\_invocation*.

**value**

```

Ar_has_impl_request : Wf_Design_Renaming → Bool
Ar_has_impl_request(dr) ≡
  self_invocation(Adapter, Request, SpecificRequest, dr)

```

The above specifies that the `SpecificRequest` methods which the `Adapter` class inherits from the `Adaptee` class are used by the `Request` methods in the `Adapter` class. In fact to conform to the Adapter pattern this should be the only use of the `SpecificRequest` methods – the inheritance relation between the `Adaptee` and `Adapter` classes is introduced in the pattern simply to make the `SpecificRequest` methods in the `Adaptee` class accessible to the `Client` through the interface of the `Target` class. The `SpecificRequest` methods, and indeed any other methods in the `Adaptee` class, although inherited by the `Adapter` class, should not be used outside that class. They thus belong to the *private* rather than the public interface of the class.

The function *has\_private\_interface\_by\_inh* expresses this property by requiring that no class can have a method including an invocation to a private method except the `Adapter` class.

**value**

```

Ar_has_private_sp_rqst : Wf_Design_Renaming → Bool
Ar_has_private_sp_rqst(dr) ≡
  has_private_interface_by_inh(Adapter, Adaptee, dr)

```

The final property describes the role of the `Client` in the pattern. The `Request` methods in the `Target` class correspond to the interface that the `Client` needs, so the interactions of the `Client` involve the invocation of these methods in some way. In the pattern structure in Figure 4, this interaction is depicted as the association relation between the two classes, though in fact the relation could be either an association or an aggregation depending on the context in which the pattern is applied. This property is formalised using the function *has\_assoc\_aggr\_reftype*.

In addition we specify that the interaction is in fact through the **Request** methods in the **Target** class. This is done using the function *use\_interface*, which requires that the **Client** class includes some method or methods, each of which contains at least one invocation to one of the **Request** methods in the **Target** class. Finally, the **Client** should not invoke the **SpecificRequest** operations of the **Adaptee** class directly – if this is possible the **Adapter** class and hence the whole pattern is unnecessary. This is specified using the function *not\_use\_interface*. Note that this does not rule out the possibility that the **Client** might be able to invoke some operations of the **Adaptee** class directly, though any such methods would not be **SpecificRequest** methods.

**value**

```

adapter_client : Wf_Design_Renaming → Bool
adapter_client(dr) ≡
  has_assoc_aggr_retype(Client, Target, AssAggr, G.one, dr) ∧
  use_interface(Client, Target, Request, dr) ∧
  not_use_interface(Client, Adaptee, SpecificRequest, dr)

```

Combining all these properties together we obtain the following specification of the function *is\_class\_adapter\_pattern* which determines whether a particular design corresponds to the Class Adapter pattern under a given renaming:

**value**

```

is_class_adapter_pattern : Wf_Design_Renaming → Bool
is_class_adapter_pattern(dr) ≡
  one_target_in_hierarchy(dr) ∧
  exists_adapter(dr) ∧
  one_adaptee_in_hierarchy(dr) ∧
  is_concrete_adaptee(dr) ∧
  adapter_client(dr) ∧
  Ar_has_two_parents(dr) ∧
  CT_is_concrete_target(dr) ∧
  subclassAdaptee_is_concrete(dr) ∧
  Ae_has_specific_request(dr) ∧
  T_has_defined_request(dr) ∧
  Ar_has_impl_request(dr) ∧ Ar_has_private_sp_rqst(dr)

```

### 2.2.2 Formal specification of the Object Adapter pattern

The basic entities appearing in the modified structure of the Object Adapter pattern shown in Figure 5 are again defined as RSL constants. Most of these are the same as those used in the specification of the Class Adapter pattern above, the differences being that the class name **SubclassAdaptee** is omitted while the variable name **adaptee** is added.

**value**

Target, ConcreteTarget, Adapter, Adaptee, Client : G.Class\_Name,

adaptee : G.Variable\_Name,

Request, SpecificRequest : G.Method\_Name

Many of the properties of the pattern entities are also the same as those of the same entities in the Class Adapter pattern, so we concentrate here on describing the differences.

Beginning again with the hierarchy of classes consisting of the abstract class **Target** together with its concrete subclasses **ConcreteTarget** and **Adapter**, the basic properties are the same as those in the Class Adapter pattern except that we do not need to explicitly exclude the role **SubclassAdaptee** from the hierarchy because this role does not occur in the Object Adapter pattern. The specification is therefore exactly the same as that for the Class Adapter except that the role **SubclassAdaptee** is omitted from the arguments of the *hierarchy* function:

**value**

one\_target\_in\_hierarchy : Wf\_Design\_Renaming → **Bool**

one\_target\_in\_hierarchy(dr) ≡

hierarchy(Target, {Adapter, ConcreteTarget}, {Client, Adaptee}, dr)

The **Adaptee** class does not belong to a hierarchy in the Object Adapter pattern so we must specify its properties differently. However, we can again assume without any loss of generality that the class is unique – if a **Client** needs to use **SpecificRequest** methods from two different **Adaptee** classes we can consider this as corresponding to two distinct instances of the Object Adapter pattern which happen to have the same **Client** and **Target** classes. We therefore use the function *exists\_one* to state directly that there is only one class in the design playing the **Adaptee** role.

**value**

exists\_one\_adaptee : Wf\_Design\_Renaming → **Bool**

exists\_one\_adaptee(dr) ≡ exists\_one(Adaptee, dr)

The property that there must be at least one class playing the **Adapter** role is shared with the Class Adapter pattern and so its specification is identical:

**value**

exists\_adapter : Wf\_Design\_Renaming → **Bool**

exists\_adapter(dr) ≡ exists\_role(Adapter, dr)

Also as before, both the **Adapter** and **ConcreteTarget** classes are concrete subclasses of the abstract **Target** class. These two properties, each of which is specified using the function *is\_concrete* (cf. the functions *CT\_is\_concrete\_target* and *Ar\_has\_two\_parents* in the specification of the Class Adapter pattern above), are combined in the function *are\_concrete\_target*:

```

value
are_concrete_target : Wf_Design_Renaming → Bool
are_concrete_target(dr) ≡
  is_concrete(Target, Adapter, dr) ∧
  is_concrete(Target, ConcreteTarget, dr)

```

The **Adaptee** class is again concrete, and the properties of the **Request** methods in the **Target** class and of the **SpecificRequest** methods in the **Adaptee** class are also the same as in the Class Adapter pattern. The specifications of these properties are therefore again given by the functions *is\_concrete\_adaptee*, *T\_has\_defined\_request* and *Ae\_has\_specific\_request*.

```

value
is_concrete_adaptee : Wf_Design_Renaming → Bool
is_concrete_adaptee(dr) ≡ is_concrete_class(Adaptee, dr),

T_has_defined_request : Wf_Design_Renaming → Bool
T_has_defined_request(dr) ≡
  has_def_method(Target, Request, dr) ∧
  has_all_def_method(Target, Request, dr),

Ae_has_specific_request : Wf_Design_Renaming → Bool
Ae_has_specific_request(dr) ≡
  has_impl_method(Adaptee, SpecificRequest, dr) ∧
  has_all_impl_method(Adaptee, SpecificRequest, dr)

```

The main difference between the Object Adapter and the Class Adapter patterns comes in the body of the **Request** method: in the Class Adapter this method involves a self-invocation of an inherited **SpecificRequest** method, whereas in the Object Adapter the invocation is to the state variable **adaptee** which represents the reference to an object belonging to the **Adaptee** class as depicted in the association relation with this class (see Figure 5). We begin by specifying the properties of this variable and relation.

Each **Adapter** class adapts the interface of a single **Adaptee**, so both the **adaptee** state variable and the association relation linking the **Adapter** and the **Adaptee** classes are unique. The existence and uniqueness of the state variable is specified using the function *store\_unique\_vble*, while the functions *has\_assoc\_aggr\_var\_ren* and *has\_unique\_assoc\_aggr\_relation* define the properties of the relation: the first is similar to the function *has\_assoc\_aggr\_reltype* used in the specification of the

function *adapter\_client* in Section 2.2.1 and also below except that it also specifies the name of the relation, in this case *adaptee*; the second function ensures that there is only one association or aggregation relation between a given pair of classes.

**value**

```
store_unique_adaptee : Wf_Design_Renaming → Bool
store_unique_adaptee(dr) ≡ store_unique_vble(Adapter, adaptee, dr),

adapter_relation : Wf_Design_Renaming → Bool
adapter_relation(dr) ≡
  has_assoc_aggr_var_ren(Adapter, Adaptee, Association, adaptee, G.one, dr) ∧
  has_unique_assoc_aggr_relation(Adapter, Adaptee, dr)
```

The specification of the **Request** method in the **Adapter** class is then simply that its body consists of a single invocation to the *adaptee* state variable of a **SpecificRequest** method – again the existence and concreteness of the **Request** method follow automatically from the fact that the **Adapter** class is a concrete subclass of the **Target** class. This is specified using the function *deleg\_with\_var*.

**value**

```
Ar_has_implemented_request : Wf_Design_Renaming → Bool
Ar_has_implemented_request(dr) ≡
  deleg_with_var(Adapter, Request, adaptee, Adaptee, SpecificRequest, dr)
```

Finally, the **Client** class plays precisely the same role in both forms of the pattern so its properties are once again specified using the function *adapter\_client*:

**value**

```
adapter_client : Wf_Design_Renaming → Bool
adapter_client(dr) ≡
  has_assoc_aggr_retype(Client, Target, AssAggr, G.one, dr) ∧
  use_interface(Client, Target, Request, dr) ∧
  not_use_interface(Client, Adaptee, SpecificRequest, dr)
```

Collecting these properties together gives the following function which defines whether or not a design matches the Object Adapter pattern:

**value**

```
is_object_adapter_pattern : Wf_Design_Renaming → Bool
```

$$\begin{aligned}
\text{is\_object\_adapter\_pattern}(\text{dr}) \equiv & \\
& \text{one\_target\_in\_hierarchy}(\text{dr}) \wedge \\
& \text{exists\_adapter}(\text{dr}) \wedge \\
& \text{exists\_one\_adaptee}(\text{dr}) \wedge \\
& \text{is\_concrete\_adaptee}(\text{dr}) \wedge \\
& \text{exists\_concrete\_target}(\text{dr}) \wedge \\
& \text{adapter\_client}(\text{dr}) \wedge \\
& \text{store\_unique\_adaptee}(\text{dr}) \wedge \\
& \text{adapter\_relation}(\text{dr}) \wedge \\
& \text{are\_concrete\_target}(\text{dr}) \wedge \\
& \text{Ae\_has\_specific\_request}(\text{dr}) \wedge \\
& \text{T\_has\_defined\_request}(\text{dr}) \wedge \\
& \text{Ar\_has\_implemented\_request}(\text{dr})
\end{aligned}$$

### 3 The Bridge Pattern

When an abstraction can have one of several possible implementations, the usual way to accommodate them is to use inheritance. An abstract class defines the interface to the abstraction, and concrete subclasses implement it in different ways. But this approach is not always flexible enough. Inheritance binds an implementation to the abstraction permanently, which makes it difficult to modify, extend, and reuse abstractions and implementations independently [7].

A better approach is to separate the abstraction from the implementation and establish a binding between them. Both abstraction and implementation can then vary independently. The binding between the abstraction and the implementation thus effectively acts as a *bridge*, because it bridges the abstraction and the implementation.

The Bridge pattern addresses these problems. The essential elements which define it, presented in the consistent format used in [7], are introduced first in Section 3.1. Then these properties are analysed and formalised in Section 3.2.

#### 3.1 Properties of the Bridge Pattern

##### Intent

Decouple an abstraction from its implementation so that the two can vary independently.

##### Structure

Figure 6 shows the structure of the Bridge Pattern.

##### Participants

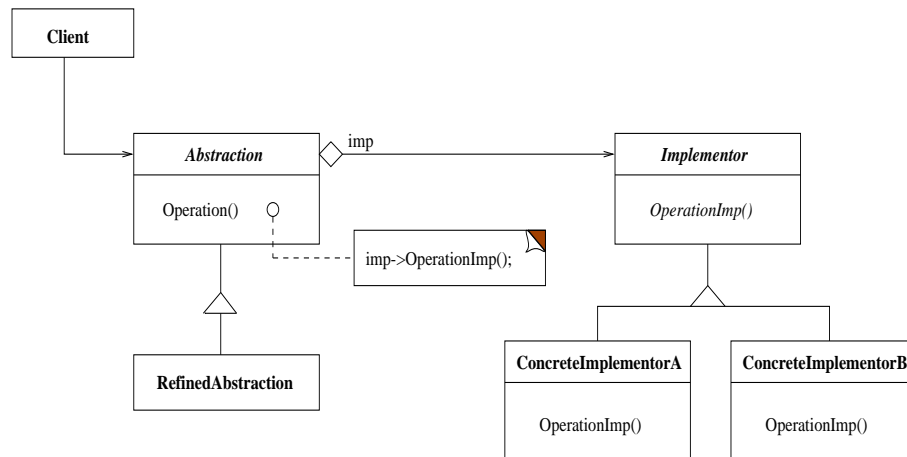


Figure 6: Bridge Pattern Structure

- **Abstraction**
  - defines the abstraction's interface.
  - maintains a reference to an object of type Implementor.
- **Refined Abstraction**
  - extends the interface defined by Abstraction.
- **Implementor**
  - defines the interface for implementation classes. This interface doesn't have to correspond exactly to Abstraction's interface; in fact the two interfaces can be quite different. Typically the Implementor interface provides only primitive operations, and Abstraction defines higher-level operations based on these primitives.
- **Concrete Implementor**
  - implements the Implementor interface and defines its concrete implementation.

### Collaborations

- Abstraction forwards client requests to its Implementor object.

## 3.2 Formalising the Bridge Pattern

From the structure of the Bridge pattern shown in Figure 6, the names of the classes in the pattern together with the names of their methods and state variables, which together describe their behaviours and properties, are defined as RSL constants.

### value

Abstraction, Implementor, RefinedAbstraction, ConcreteImplementor,



```

Client : G.Class_Name,

Operation, OperationImp : G.Method_Name,

imp : G.Variable_Name

```

The structure in Figure 6 shows linked inheritance hierarchies of classes in which the *root* classes are the classes **Abstraction** and **Implementor** respectively. Both of these root classes should be unique because their purpose is to factor out the common interface of their subclasses into abstract classes. In addition, the particular classes playing these roles must not play any other role in the pattern (for example the **Abstraction** class cannot also play the **Implementor** or the **ConcreteImplementor** role) because this would destroy the separation between the abstraction and the implementation which is the essence of the pattern.

For the same reason, the subclasses in the two hierarchies may play only a single role in the pattern, so that subclasses of the **Abstraction** class can play only the **RefinedAbstraction** role and subclasses of the **Implementor** class can only play the **ConcreteImplementor** role. However, there can be more than one subclass playing each of these roles. Indeed in general it is extremely likely that this would be the case: while a single subclass is not impossible, there is in fact nothing to be gained by defining an abstract superclass in such a situation and we could instead combine the two roles. We do not deal with this possibility here, however, preferring to consider it as a *variant* of the pattern. In fact there are many possible variants of patterns, many of which are discussed in the literature when authors wish to represent refinements of or extensions to the basic patterns. Specification of such variants of the structural patterns will be the subject of future work.

The above properties of the two class hierarchies are analogous to the properties of the hierarchies in the Adapter pattern discussed in Section 2.2.1, so they are also described using the function *hierarchy*, here instantiated with the appropriate roles from the Bridge pattern.

**value**

```

one_abstraction_in_hierarchy : Wf_Design_Renaming → Bool
one_abstraction_in_hierarchy(dr) ≡
  hierarchy
    (Abstraction, {RefinedAbstraction},
     {Client, Implementor, ConcreteImplementor}, dr),

one_implementor_in_hierarchy : Wf_Design_Renaming → Bool
one_implementor_in_hierarchy(dr) ≡
  hierarchy
    (Implementor, {ConcreteImplementor},
     {Client, Abstraction, RefinedAbstraction}, dr)

```

The root class of each hierarchy should be an abstract class. In the case of the `Implementor` class, this property follows automatically from the fact that the class contains the abstract `OperationImp` method (cf. the discussion of the `Target` class in Section 2.2.1). For the `Abstraction` class, however, we must specify this property explicitly. This is done using the function *is\_abstract\_class* which is analogous to the function *is\_concrete\_class* used to specify that the `Adaptee` class in the Class Adapter pattern is concrete (see Section 2.2.1).

**value**

```
abstraction_is_abstract_class : Wf_Design_Renaming → Bool
abstraction_is_abstract_class(dr) ≡ is_abstract_class(Abstraction, dr)
```

Each class playing the `ConcreteImplementor` role describes a different implementation for abstractions. At least one class in a design must provide such an implementation otherwise it does not make sense to apply the Bridge pattern. In situations where there is only one implementation, creating an abstract `Implementor` class with one `ConcreteImplementor` subclass is not necessary – in this case every abstraction would be using the same implementation so we could omit the abstract superclass `Implementor`. Nevertheless, the separation is not incorrect and can be useful when the domain represents a variable context in which another implementation can be added later. We therefore simply specify that there must be at least one `ConcreteImplementor` and, by similar arguments, at least one `RefinedAbstraction` class. This is done using the function *exists\_role* exactly as in the specification of the existence of the `Adapter` class in the Class Adapter pattern (see Section 2.2.1).

**value**

```
exists_concrete_implementor : Wf_Design_Renaming → Bool
exists_concrete_implementor(dr) ≡ exists_role(ConcreteImplementor, dr),

exists_refined_abstraction : Wf_Design_Renaming → Bool
exists_refined_abstraction(dr) ≡ exists_role(RefinedAbstraction, dr)
```

The classes playing the `RefinedAbstraction` role are concrete subclasses of the `Abstraction` class, which is specified using the function *is\_concrete* as in the case of the `ConcreteTarget` and `SubclassAdaptee` classes in the Class Adapter pattern described in Section 2.2.1. In addition, they inherit at least one implemented `Operation` method from the `Abstraction` class which represents the common basic behaviour of an abstraction. The `Abstraction` class is abstract even though its `Operation` methods are concrete because there are usually many different abstractions which are represented as subclasses of this abstract class. So each class playing the `RefinedAbstraction` role describes a “refined” or more specific abstraction and therefore extends the interface of its parent. We specify this using the function *extends\_interface*, which states that every class playing a given role (`RefinedAbstraction`) must either implement abstract methods it inherits from its parent classes or must extend the interface of its parent classes by adding some new methods or state variables.

**value**

```
is_concrete_refined_abstraction : Wf_Design_Renaming → Bool
is_concrete_refined_abstraction(dr) ≡
  is_concrete(Abstraction, RefinedAbstraction, dr) ∧
  extends_interface(Abstraction, RefinedAbstraction, dr)
```

Similarly the `ConcreteImplementor` classes must be concrete subclasses of the `Implementor` class, which is again specified using the function *is\_concrete*:

**value**

```
is_concrete_implementor : Wf_Design_Renaming → Bool
is_concrete_implementor(dr) ≡ is_concrete(Implementor, ConcreteImplementor, dr)
```

The structure of the Bridge pattern in Figure 6 shows that the `Implementor` class contains an abstract `OperationImp` method. This in fact represents the interface for the implementation classes mentioned in the description of the pattern's participants, so in principle there can be more than one method playing this role in a design. We formalise this property using the functions *has\_def\_method* and *has\_all\_def\_method* functions exactly as in the specification of the properties of the `Request` methods in the `Target` class in the Class Adapter pattern (see Section 2.2.1).

**value**

```
Lhas_def_operationImp : Wf_Design_Renaming → Bool
Lhas_def_operationImp(dr) ≡
  has_def_method(Implementor, OperationImp, dr) ∧
  has_all_def_method(Implementor, OperationImp, dr)
```

From the structure of the pattern it is clear that the `ConcreteImplementor` classes are concrete. However, a concrete class can contain error methods as well as implemented methods (see [6]), so we need to explicitly specify that the `OperationImp` methods are actually implemented. This specification is analogous to that immediately above except that we use the functions *has\_impl\_method* and *has\_all\_impl\_method*.

**value**

```
CIhas_impl_operation : Wf_Design_Renaming → Bool
CIhas_impl_operation(dr) ≡
  has_impl_method(ConcreteImplementor, OperationImp, dr) ∧
  has_all_impl_method(ConcreteImplementor, OperationImp, dr)
```

In order to allow instances of the **RefinedAbstraction** classes to vary their implementation dynamically, they store a reference to an instance of a subclass of the **Implementor** class in their **imp** state variable, which is shown in the pattern structure as the name of the aggregation relation linking the **Abstraction** class to the **Implementor** class. Both the variable and the relation should be unique because each abstraction has a single implementation. These properties are specified using the functions *store\_unique\_vble*, *has\_assoc\_aggr\_var\_ren* and *has\_unique\_assoc\_aggr\_relation* and are entirely analogous to the properties of the relation linking the **Adapter** and the **Adaptee** classes in the Object Adapter pattern (see Section 2.2.2).

**value**

```
store_unique_imp : Wf_Design_Renaming → Bool
store_unique_imp(dr) ≡ store_unique_vble(Abstraction, imp, dr),

bridge_relation : Wf_Design_Renaming → Bool
bridge_relation(dr) ≡
  has_assoc_aggr_var_ren(Abstraction, Implementor, Aggregation, imp, G.one, dr) ∧
  has_unique_assoc_aggr_relation(Abstraction, Implementor, dr)
```

The **Abstraction** class contains at least one implemented **Operation** method, the objective of which is to forward the **Client** requests by invoking the **OperationImp** methods in the **Implementor** class on the **imp** state variable. Neither the result nor the parameters of the invocation are fixed by the pattern because they may be different in different methods. This property is specified using the function *deleg\_with\_var* exactly as in the specification of the **Request** method in the **Adapter** class in the Object Adapter pattern.

**value**

```
A_has_impl_operation : Wf_Design_Renaming → Bool
A_has_impl_operation(dr) ≡
  has_impl_method(Abstraction, Operation, dr) ∧
  deleg_with_var(Abstraction, Operation, imp, Implementor, OperationImp, dr)
```

Finally, we specify the role of the **Client** in the pattern. According to the Collaborations, the **Client** sends requests to the **Abstraction** class which are then forwarded to the **Implementor** class as described above. To do this, the **Client** must invoke the **Operation** method in the **Abstraction** class, which is represented in the pattern structure by the relation linking the **Client** and **Abstraction** classes which could be either an association or an aggregation relation depending of the context in which the pattern is applied. In addition, in order to preserve the intent of the Bridge pattern the **Implementor** interface should not be directly accessed by the **Client** class so that the **Client** does not need to know about the **Implementor** or its **ConcreteImplementor** subclasses. This means that the **Client** does not need to be changed every time the **ConcreteImplementor** is changed, thus making the **Client** class reusable by changing its implementation dynamically rather than statically committing it to a particular implementation.

The interaction of the Client with the **Abstraction** class is therefore analogous to the interaction of the Client with the **Target** class in the Adapter pattern, so the specification of these properties again uses the functions *has\_assoc\_aggr\_retype* and *use\_interface*. However, in the Bridge pattern the Client may not directly access the **Implementor** class whereas in the Adapter pattern it is only forbidden to access the operations playing the **SpecificRequest** role. We therefore use the function *not\_related\_classes* instead of the function *not\_use\_interface* to capture this stronger property here.

**value**

```
bridge_client : Wf_Design_Renaming → Bool
bridge_client(dr) ≡
  has_assoc_aggr_retype(Client, Abstraction, AssAggr, G.one, dr) ∧
  use_interface(Client, Abstraction, Operation, dr) ∧
  not_related_classes(Client, Implementor, dr)
```

Collecting all these properties together yields the function *is\_bridge\_pattern* which defines whether or not a design matches the Bridge pattern:

**value**

```
is_bridge_pattern : Wf_Design_Renaming → Bool
is_bridge_pattern(dr) ≡
  one_abstraction_in_hierarchy(dr) ∧
  one_implementor_in_hierarchy(dr) ∧
  bridge_client(dr) ∧
  exists_concrete_implementor(dr) ∧
  exists_refined_abstraction(dr) ∧
  abstraction_is_abstract_class(dr) ∧
  is_concrete_refined_abstraction(dr) ∧
  is_concrete_implementor(dr) ∧
  store_unique_imp(dr) ∧
  bridge_relation(dr) ∧
  A_has_impl_operation(dr) ∧
  I_has_def_operationImp(dr) ∧
  CI_has_impl_operation(dr)
```

## 4 The Composite Pattern

The Composite pattern belongs to the group of patterns which are based on some form of recursive structure, which is one aspect identified in [10] as offering another way of classifying the

GoF design patterns. It describes how to represent hierarchically structured information using a technique called recursive composition [7] and supports the representation of any potentially complex hierarchical structure. In addition, the Composite structure allows clients to treat the composite object and the component objects uniformly by offering a single interface to both. It can thus help to reduce complexity by allowing many objects to be treated as a single object [9].

We start by introducing the properties of the pattern as described in [7], then we analyse these properties and present our formalisation of them in our generic formal model [6].

## 4.1 Properties of the Composite Pattern

### Intent

Compose objects into tree structures to represent part-whole hierarchies. Composite lets clients treat individual objects and compositions of objects uniformly.

### Structure

Figure 7 shows the structure of the Composite Pattern, and Figure 8 shows a typical Composite object structure.

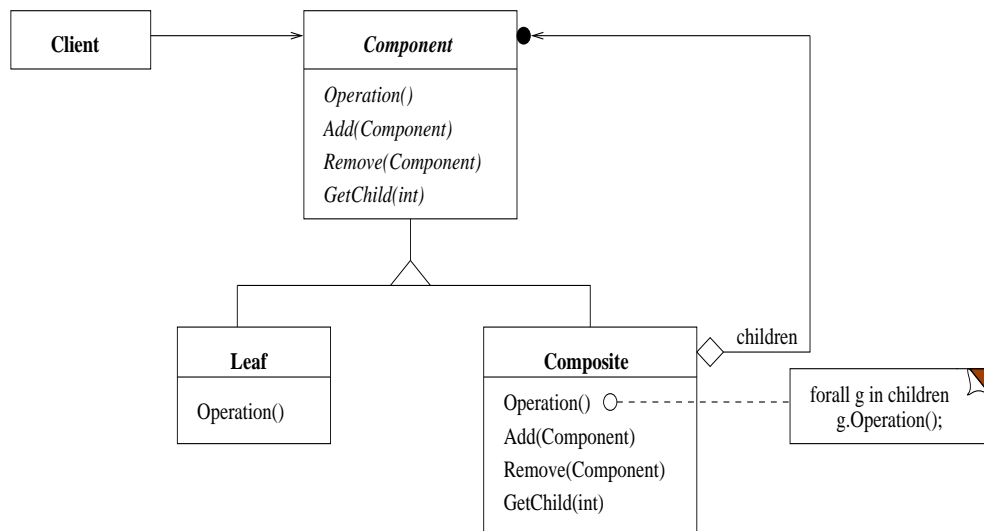


Figure 7: Composite Structure

### Participants

- **Component**
  - declares the interface for objects in the composition.
  - implements default behaviour for the interface common to all classes, as appropriate.

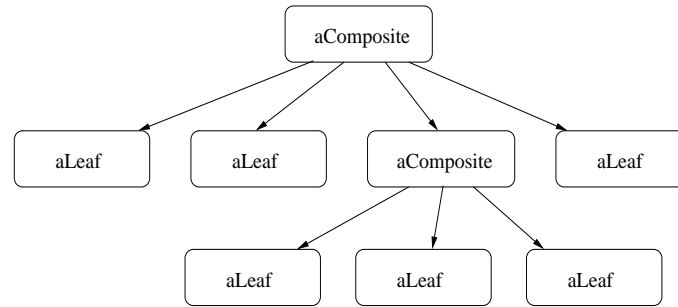


Figure 8: Composite Object Structure

- declares an interface for accessing and managing its child components.
- Leaf
  - represents leaf objects in the composition. A leaf has no children.
  - defines behaviour for primitive objects in the composition.
- Composite
  - defines behaviour for components having children.
  - stores child components.
  - implements child-related operations in the Component interface.
- Client
  - manipulates objects in the composition through the Component interface.

### Collaborations

- Clients use the component class interface to interact with objects in the composite structure. If the recipient is a leaf, then the request is handled directly. If the recipient is a Composite, then it usually forwards requests to its child components, possibly performing additional operations before and/or after forwarding.

## 4.2 Formalising the Composite Pattern

We begin as usual by defining RSL constants representing the roles of the classes, methods and variables which participate in the pattern:

### value

Component, Composite, Leaf, Client : G.Class\_Name,

Operation, Add, Remove, GetChild : G.Method\_Name,

children, component, inte : G.Variable\_Name

As shown in Figure 7, the Composite structure basically comprises a single class hierarchy in which the root (the **Component** class) provides a common interface to both the **Leaf** and the **Composite** classes, where the **Leaf** class represents the individual objects and the **Composite** class represents the composite objects. The **Client** interacts solely through the interface of the abstract **Component** class, thus enabling it to treat objects of the **Leaf** and **Composite** classes uniformly. Clients can thus deal with these objects without knowing whether they are composite or primitive.

Without loss of generality we can assume that there is only one class playing the **Component** role – this class declares a common interface for its subclasses, so that a second such class would require a second hierarchy of subclasses and this can be considered as a second and separate instance of the pattern (possibly with the same client). In addition, it is not possible for a single class to play both the **Leaf** and the **Composite** roles because these represent different types of object (primitive and composite respectively). The properties of this hierarchy are therefore once again analogous to those of the hierarchy of **Target** classes in the Class Adapter pattern described in Section 2.2.1 so are similarly specified using an appropriate instantiation of the function *hierarchy*. Note that again we do not need to specify explicitly that the **Component** class is abstract because this again follows from the fact that it contains abstract methods.

**value**

```
one_component_in_hierarchy : Wf_Design_Renaming → Bool
one_component_in_hierarchy(dr) ≡
  hierarchy(Component, {Leaf, Composite}, {Client}, dr)
```

The **Leaf** and the **Composite** classes represent the actual objects, primitive and composite respectively, with which the **Client** interacts. These classes must therefore be concrete subclasses of the **Component** class. In addition, there must be at least one class in the design playing each of these roles – at least one **Leaf** class is necessary in order for the recursion implicit in the **Composite** class to be finite, and the **Composite** class is the essential part of the pattern so without it there is no mechanism for structuring objects and hence no need for the pattern. In fact there may be many classes in a design which play each of these roles. In the case of the **Leaf** role this is clear since structures can obviously be composed of many different primitive objects. An example of a design in which more than one class plays the **Composite** role can be seen in the sample code for the Composite pattern presented in [7], where there are three classes playing this role. These properties of the **Leaf** and **Composite** classes are therefore analogous to the corresponding properties of the **Adapter** class in the Class Adapter pattern (see Section 2.2.1) and so they are similarly specified using the functions *exists\_role* and *is\_concrete*.

**value**

```
exist_composite : Wf_Design_Renaming → Bool
exist_composite(dr) ≡ exists_role(Composite, dr),

exist_leaf : Wf_Design_Renaming → Bool
```



```

exist_leaf(dr)  $\equiv$  exists_role(Leaf, dr),

are_concrete_component : Wf_Design_Renaming  $\rightarrow$  Bool
are_concrete_component(dr)  $\equiv$ 
  is_concrete(Component, Composite, dr)  $\wedge$ 
  is_concrete(Component, Leaf, dr)

```

The main feature of the Composite pattern is the recursive composition of objects which allows object hierarchies to be built dynamically [7, 3]. In order to achieve this recursion, objects of the **Composite** class store in their state a collection of objects, each of which may be another composite object or a leaf object and which are therefore in general of the **Component** type. This collection of sub-objects is represented by the state variable *children*, which also corresponds to the name of the aggregation relation between the **Composite** and **Component** classes. Through this relation, which is called “recursive aggregation” in [13, 3], it is possible to compose objects into tree structures at run-time (see, for example, Figure 8 in Section 4.2 which illustrates an object hierarchy which has been built in this way), and since the sink of the relation is the **Component** class, which is abstract, the objects will actually belong to the **Leaf** or the **Composite** class.

Of course the relation is one-many since a composite object can, and in general will be constructed from more than one sub-object. It is also unique since the role of the **Component** class is simply to provide a common, abstract interface to both the **Leaf** and the **Composite** classes and the relation between the **Composite** and **Component** classes exists solely to implement the recursive composition.

These properties are therefore analogous to the properties of the *adaptee* state variable in the **Adapter** class in the Object Adapter pattern and its corresponding relation with the **Adaptee** class (see Section 2.2.2), so they are similarly specified using the functions *store\_unique\_vble*, *has\_assoc\_aggr\_var\_ren* and *has\_unique\_assoc\_aggr\_relation*.

**value**

```

store_unique_children : Wf_Design_Renaming  $\rightarrow$  Bool
store_unique_children(dr)  $\equiv$  store_unique_vble(Composite, children, dr),

composite_relation : Wf_Design_Renaming  $\rightarrow$  Bool
composite_relation(dr)  $\equiv$ 
  has_assoc_aggr_var_ren(Composite, Component, Aggregation, children, G.many, dr)  $\wedge$ 
  has_unique_assoc_aggr_relation(Composite, Component, dr)

```

As explained above, the **Component** class defines the interface which is common to all its sub-classes, both **Leaf** classes and **Composite** classes. This interface basically consists of two parts: what might be termed the “true” interface, which consists of those operations that clients can

invoke on the component objects and which is represented by the method role **Operation**; and child management operations for manipulating the structure of **Composite** objects, specifically for adding a new component to a composite object (the **Add** operation), for removing a component from a composite object (the **Remove** operation), and for returning one of the components of a composite object (the **GetChild** operation).

As far as the interface represented by the **Operation** method is concerned, these operations in the design describe the main functionality of the primitive (leaf) objects. Therefore there will in general be many such methods in each **Leaf** class, and hence also in the **Component** and **Composite** classes since these must share the same interface.

In the **Component** class each of these methods will be abstract since they must be implemented differently in the **Leaf** and the **Composite** classes – in the **Leaf** classes they perform actual operations on the primitive objects whereas in the **Composite** classes they are simply forwarded to the sub-objects. At this level, therefore, we can only specify that at least one such method must exist and that all such methods must be abstract. These properties are thus analogous to those of the **Request** method in the **Target** class in the Class Adapter pattern and so the specification is again given in terms of the functions *has\_def\_method* and *has\_all\_def\_method*.

**value**

```
Ct_has_def_operation : Wf_Design_Renaming → Bool
Ct_has_def_operation(dr) ≡
  has_def_method(Component, Operation, dr) ∧
  has_all_def_method(Component, Operation, dr)
```

On the other hand, the methods must all be implemented in all the **Leaf** classes and in all the **Composite** classes. In the **Leaf** classes the methods are handled directly, that is the **Leaf** class implements the **Operation** methods so as to provide the appropriate functionality for the particular type of sub-object it represents. These implementations will of course usually be different in different **Leaf** classes, so we cannot specify anything about the implementations. In the **Composite** classes, however, the functionality of the methods is precisely determined – each **Operation** method simply invokes the same method on all the object's sub-objects as represented by the state variable *children*. Since the general model in [6] automatically interprets an invocation of a method on a collection of objects as individual invocations of the same method on each object in the collection, this is analogous to the behaviour of the **Request** method in the Object Adapter pattern and so is similarly specified using the function *deleg\_with\_var*. Note that this function also implies that the methods are implemented so we do not need to state this explicitly.

In addition, the **Component** class should represent the whole of the interface available to clients, which means that neither the **Leaf** nor the **Composite** classes can introduce new **Operation** methods. This property is specified using the function *no\_adds\_method\_role* which is one of the auxiliary functions which are defined in [2] as extensions to the basic model discussed in [6].

**value**

```

Ce_has_impl_operation : Wf_Design_Renaming → Bool
Ce_has_impl_operation(dr) ≡
  deleg_with_var(Composite, Operation, children, Component, Operation, dr) ∧
  no_adds_method_role(Composite, Component, Operation, dr),

L_has_impl_operation : Wf_Design_Renaming → Bool
L_has_impl_operation(dr) ≡
  has_all_impl_method(Leaf, Operation, dr) ∧
  no_adds_method_role(Leaf, Component, Operation, dr)

```

Turning now to the child management operations **Add**, **Remove**, and **GetChild**, each of these should also be abstract at the level of the **Component** class. However, these methods perform specific operations on the **children** of a composite object so only one of each method is required. We specify this uniqueness, together with the functionality of each method, at the level of the **Composite** class – if a method is unique in some class then there cannot be more than one such method in any superclass of that class.

The **Composite** class implements the **Add** and **Remove** methods by invoking the *collectionadd* and *collectionremove* methods respectively on the state variable **children**. These methods, which represent abstractions of the range of operations that are found in object-oriented programming languages for manipulating different types of collections of objects, respectively add and remove a given object (represented by their single parameter) from a given generic (i.e. of any type: set, list, etc.) collection of objects (represented by the receiver of the message) and are built into the basic model described in [6], their names being reserved method names. The bodies of the **Add** and **Remove** methods thus consist of a single invocation to the state variable **children** of the appropriate collection manipulation method, the parameter **component** of the **Add** or **Remove** method being passed directly as a parameter to the invocation. This behaviour is represented by the function *deleg\_with\_var\_coll\_aparam\_ren* from [6], and this, together with the function *unique\_method* which ensures the uniqueness of the method, forms the specification of the **Add** and **Remove** methods at the level of the **Composite** class.

**value**

```

Ce_has_impl_add : Wf_Design_Renaming → Bool
Ce_has_impl_add(dr) ≡
  unique_method(Composite, Add, dr) ∧
  deleg_with_var_coll_aparam_ren
    (Composite, Add, children, G.collectionadd, component, dr),

Ce_has_impl_remove : Wf_Design_Renaming → Bool
Ce_has_impl_remove(dr) ≡
  unique_method(Composite, Remove, dr) ∧
  deleg_with_var_coll_aparam_ren
    (Composite, Remove, children, G.collectionremove, component, dr)

```

The functionality of the `GetChild` method in the `Composite` class is similarly described using an invocation to a generic collection manipulation method, in this case `collectionelement` which returns the object at a given “location” in a generic collection. Again, the location is represented by the method’s single parameter and the collection is the receiver of the message. The body of the `GetChild` method thus similarly consists of a single invocation to the `children` state variable of the `collectionelement` method, the parameter `inte`<sup>2</sup> of the `GetChild` method being passed as a parameter to the invocation. However, in this case the method also returns a result, which is in fact the result of the invocation of the `collectionelement` method, and according to the general model in [6] this means that the result returned by the invocation of the `collectionelement` method must be assigned to a local variable in the variable change of the `GetChild` method and this local variable is then returned as the result of the method. These properties are embodied in the generic function *res\_local\_var\_change\_inv\_aparam\_ren* from [6], and this is therefore used together with the function *unique\_method* once more to give the following specification of the `GetChild` method in the `Composite` class.

**value**

```
Ce_has_impl_get_child : Wf_Design_Renaming → Bool
Ce_has_impl_get_child(dr) ≡
  unique_method(Composite, GetChild, dr) ∧
  res_local_var_change_inv_aparam_ren
    (Composite, GetChild, children, G.collectionelement, inte, dr)
```

In the `Component` class each of these child management operations must be abstract, and since we have already specified that the methods are unique in the `Composite` classes it is sufficient to specify, using the function *has\_defined\_method*, that there is least one abstract method playing each of the three roles in the `Component` class. In addition, we specify the properties of the parameters and the result of each method at this level.

In the case of the `Add` and `Remove` methods, a single parameter, which we denote by the variable `component`, is required and this represents an object of type `Component`. This is specified using the function *one\_image\_ren\_pars\_in\_design* from [6]. Furthermore, neither of these methods returns a result since their purpose is to modify the collection of sub-objects of a composite object, that is to update the `children` variable in the `Composite` class. The function *has\_method\_without\_res* from [6] captures this property. The specifications of the properties of the `Add` and `Remove` methods at the level of the `Component` class are thus:

**value**

```
Ct_has_def_add : Wf_Design_Renaming → Bool
Ct_has_def_add(dr) ≡
  has_def_method(Component, Add, dr) ∧
```

---

<sup>2</sup>In the pattern structure the original name is `int` but we have modified it because this is a reserved keyword in RSL.

```

one_image_ren_pars_in_design(Component, Add, Component, component, dr) ∧
has_method_without_res(Component, Add, dr),

```

```

Ct_has_def_remove : Wf_Design_Renaming → Bool
Ct_has_def_remove(dr) ≡
  has_def_method(Component, Remove, dr) ∧
  one_image_ren_pars_in_design(Component, Remove, Component, component, dr) ∧
  has_method_without_res(Component, Remove, dr)

```

The `GetChild` method also has a parameter, `inte`, but in this case we do not specify the type of the parameter so we use the function *images\_ren\_one\_var\_par\_in\_design* to specify its properties instead of *one\_image\_ren\_pars\_in\_design*. In addition, the method has a result, so the specification involves the function *has\_method\_with\_res* instead of *has\_method\_without\_res*. The specification is thus:

**value**

```

Ct_has_def_get_child : Wf_Design_Renaming → Bool
Ct_has_def_get_child(dr) ≡
  has_def_method(Component, GetChild, dr) ∧
  images_ren_one_var_par_in_design(Component, GetChild, inte, dr) ∧
  has_method_with_res(Component, GetChild, dr)

```

Lastly, we consider the `Leaf` classes, where in fact the child management operations do not make sense at all although they are in principle available because they are inherited from the `Component` class. However, as explained in [7], there is a trade off between transparency and security when defining the child management operations: defining them in the `Component` class offers transparency but not security because the whole interface available to clients is then defined at this level (transparency) but security is compromised since it is possible for clients to invoke the child management operations on `Leaf` classes where they are meaningless; on the other hand, defining them solely in the `Composite` class offers security but not transparency because they are then not available in `Leaf` classes and clients have to interact directly with the `Composite` class in order to invoke them. The structure of the Composite pattern shown in Figure 7 in fact corresponds to the first of these alternatives, that is it emphasises transparency over security, and our specification follows this<sup>3</sup>. However, the general model in [6] allows us to specify explicitly that the child management operations should not be implemented in the `Leaf` classes: they are instead declared to be *error* methods in these classes using the function *has\_error\_method*. We also specify that the `Leaf` classes do not introduce new methods playing any of the child management method roles by requiring that each of these methods is unique as in the `Composite` classes. The specification of these properties is then as follows:

**value**

---

<sup>3</sup>But we could of course write a similar specification for the other alternative.

```

L_has_add_error : Wf_Design_Renaming → Bool
L_has_add_error(dr) ≡
  unique_method(Leaf, Add, dr) ∧
  has_error_method(Leaf, Add, dr),

L_has_remove_error : Wf_Design_Renaming → Bool
L_has_remove_error(dr) ≡
  unique_method(Leaf, Remove, dr) ∧
  has_error_method(Leaf, Remove, dr),

L_has_get_child_error : Wf_Design_Renaming → Bool
L_has_get_child_error(dr) ≡
  unique_method(Leaf, GetChild, dr) ∧
  has_error_method(Leaf, GetChild, dr)

```

Finally, we describe the role of the **Client** in the pattern. As explained above, both the structure shown in Figure 7 and our specification are based on the form of the pattern which emphasises transparency over security. This means that the **Client** only interacts with the **Component** class, which it may do by invoking either an **Operation** method or a child management operation (though it does not necessarily invoke the child management operations itself). This interaction is represented by the relation linking the **Client** and **Component** classes, which is shown in the structure as an association relation but which may in fact be an aggregation relation in certain designs. The **Client** therefore interacts with the **Component** class in exactly the same way as the **Client** class in the Bridge pattern interacts with the **Abstraction** class (see Section 3.2) so the specification is similarly written in terms of the functions *has\_assoc\_aggr\_reltype* and *use\_interface*.

```

value
composite_client : Wf_Design_Renaming → Bool
composite_client(dr) ≡
  has_assoc_aggr_reltype(Client, Component, AssAggr, G.one, dr) ∧
  use_interface(Client, Component, Operation, dr)

```

Combining all these properties then leads to the function *is\_composite\_pattern* which defines whether or not a design matches the Composite pattern:

```

value
is_composite_pattern : Wf_Design_Renaming → Bool
is_composite_pattern(dr) ≡
  one_component_in_hierarchy(dr) ∧
  Ct_has_def_operation(dr) ∧
  Ct_has_def_add(dr) ∧

```

$$\begin{aligned}
& \text{Ct\_has\_def\_remove}(\text{dr}) \wedge \\
& \text{Ct\_has\_def\_get\_child}(\text{dr}) \wedge \\
& \text{exist\_composite}(\text{dr}) \wedge \\
& \text{store\_unique\_children}(\text{dr}) \wedge \\
& \text{Ce\_has\_impl\_operation}(\text{dr}) \wedge \\
& \text{Ce\_has\_impl\_add}(\text{dr}) \wedge \\
& \text{Ce\_has\_impl\_remove}(\text{dr}) \wedge \\
& \text{Ce\_has\_impl\_get\_child}(\text{dr}) \wedge \\
& \text{exist\_leaf}(\text{dr}) \wedge \\
& \text{L\_has\_impl\_operation}(\text{dr}) \wedge \\
& \text{L\_has\_add\_error}(\text{dr}) \wedge \\
& \text{L\_has\_remove\_error}(\text{dr}) \wedge \\
& \text{L\_has\_get\_child\_error}(\text{dr}) \wedge \\
& \text{composite\_client}(\text{dr}) \wedge \\
& \text{are\_concrete\_component}(\text{dr}) \wedge \text{composite\_relation}(\text{dr})
\end{aligned}$$

## 5 The Decorator Pattern

Sometimes the requirements of a system include the addition of responsibilities to individual objects instead of to an entire class. Perhaps the primary way of achieving this is by inheritance: inheriting a responsibility from another class can decorate every subclass instance. However, this is inflexible since the responsibility is allocated statically, which means that a client is not able to control how and when to decorate the component. Another approach, which is more flexible, is to enclose the component inside another object which takes care of adding the responsibility. In this approach the enclosing object, which is called a *decorator*, conforms to the interface of the component it decorates so that its presence is transparent to the component's clients. In addition, it forwards requests to the component and may perform additional actions (related with the new responsibility) before or after forwarding. Decorators may be recursively nested, which effectively allows an unlimited number of added responsibilities.

The Decorator pattern embodies this second approach. It lets decorators appear anywhere a component can appear so that clients generally cannot tell the difference between a decorated component and an undecorated one and thus do not depend at all on the decoration [7]. We begin by introducing the essential elements of the pattern as described in [7], then we present our analysis and formal specification of the properties of these elements.

### 5.1 Properties of the Decorator Pattern

#### Intent

Attach additional responsibilities to an object dynamically. Decorators provide a flexible alternative to subclassing for extending functionality.

### Structure

Figure 9 shows the structure of the Decorator pattern, and Figure 10 shows an object diagram of a component composed with two decorator objects which extend its functionality.

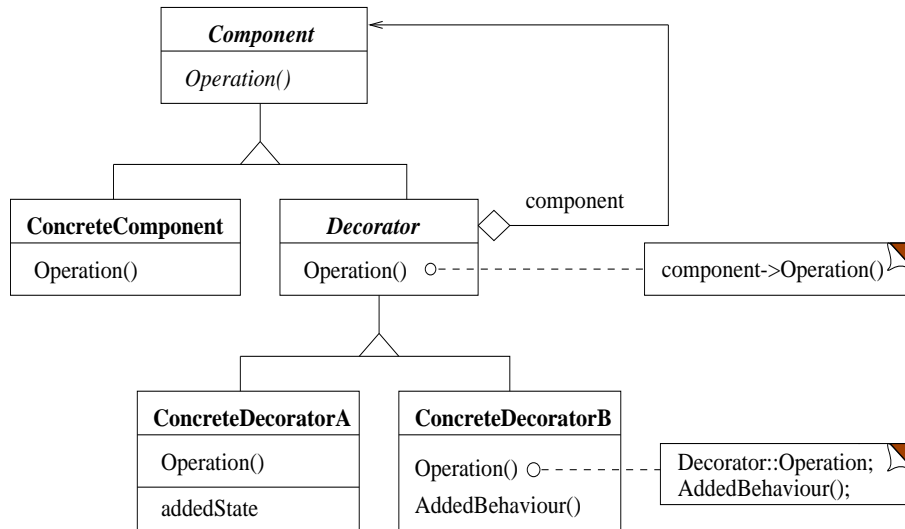


Figure 9: Decorator Pattern Structure




Figure 10: Decorator Pattern Object Structure

### Participants

- **Component**
  - defines the interface for objects that can have responsibilities added to them dynamically.
- **Concrete Component**
  - defines an object to which additional responsibilities can be attached.
- **Decorator**
  - maintains a reference to a **Component** object and defines an interface that conforms to **Component**'s interface.



- Concrete Decorator
  - adds responsibilities to the component.

### Collaborations

- Decorator forwards requests to its Component object. It may optionally perform additional operations before and after forwarding the request.

## 5.2 Formalising the Decorator Pattern

The following RSL constants represent the various entities involved in the Decorator pattern:

### value

Component, ConcreteComponent, Decorator, ConcreteDecorator, Client : G.Class\_Name,  
 Operation, AddedBehaviour : G.Method\_Name,  
 component, addedState : G.Variable\_Name

The structure of the pattern in Figure 9 shows two nested hierarchies of classes, the hierarchy of **Component** classes containing a sub-hierarchy of **Decorator** classes.

The **Decorator** class is abstract, and defines the “default” properties of all decorator objects through its **Operation** interface, which may be extended in subclasses as shown in the annotation to the method in the **ConcreteDecoratorB** class. This means that there should be only one class in the design that plays the **Decorator** role, and subclasses of this may only play the **ConcreteDecorator** role (i.e. they may not play the **ConcreteComponent** role). This hierarchy thus satisfies the same properties as the class hierarchies in the other patterns discussed above, so we can again specify its properties using the function *hierarchy*. However, in this case we must additionally specify explicitly that the **Decorator** class is abstract using the function *is\_abstract\_class*: it is a consequence of the general model in [6] that a class containing an abstract method must itself be abstract, but the **Decorator** class has a concrete interface which does not automatically determine the type of the class. The specification of the hierarchy of decorator classes is thus as follows:

### value

one\_decorator\_in\_hierarchy : Wf\_Design\_Renaming  $\rightarrow$  **Bool**  
 one\_decorator\_in\_hierarchy(dr)  $\equiv$   
 hierarchy(Decorator, {ConcreteDecorator}, {ConcreteComponent}, dr),  
  
 is\_abstract\_decorator : Wf\_Design\_Renaming  $\rightarrow$  **Bool**  
 is\_abstract\_decorator(dr)  $\equiv$  is\_abstract\_class(Decorator, dr)

Actual decorators are represented by the `ConcreteDecorator` classes, which inherit the default decorator behaviour from the `Decorator` class and extend it to match their own particular needs. The classes playing the `ConcreteDecorator` role are thus concrete subclasses of the class playing the `Decorator` role, and at least one such class must exist otherwise the design does not allow decorator objects to be created (because the `Decorator` class is abstract and so cannot be instantiated). These properties are analogous to the properties of the `ConcreteTarget` class in the Class Adapter pattern described in Section 2.2.1, so they are similarly specified using the functions *is\_concrete* and *exists\_role* respectively.

**value**

`exist_concrete_decorator : Wf_Design_Renaming → Bool`

`exist_concrete_decorator(dr) ≡ exists_role(ConcreteDecorator, dr),`

`is_concrete_decorator : Wf_Design_Renaming → Bool`

`is_concrete_decorator(dr) ≡ is_concrete(Decorator, ConcreteDecorator, dr)`

The properties of the component hierarchy are similar. The `Component` class is abstract, though in this case this follows automatically from the fact that its `Operation` method is abstract so we do not need to specify it explicitly, and we can again assume without loss of generality that there is only one class in the design that plays this role: as in the other patterns we can consider a design containing multiple `Component` classes to correspond to multiple instances of the Decorator pattern. In addition, it forms the root of the hierarchy, the leaves in which play either the `ConcreteComponent` role or the `ConcreteDecorator` role. In this case we specify that no class in the hierarchy may play the `Client` role, which in fact permits classes, but not leaf classes, to play the `Decorator` role. Finally, there must be at least one `ConcreteComponent` class otherwise the design does not permit the creation of component objects, and each such class must be a concrete subclass of the `Component` class. The specification of these properties is analogous to that of the corresponding properties of the decorator hierarchy.

**value**

`one_component_in_hierarchy : Wf_Design_Renaming → Bool`

`one_component_in_hierarchy(dr) ≡`

`hierarchy(Component, {ConcreteComponent, ConcreteDecorator}, {Client}, dr),`

`exist_concrete_component : Wf_Design_Renaming → Bool`

`exist_concrete_component(dr) ≡ exists_role(ConcreteComponent, dr),`

`is_concrete_component : Wf_Design_Renaming → Bool`

`is_concrete_component(dr) ≡ is_concrete(Component, ConcreteComponent, dr)`

In order that the decorators should be transparent to the clients, the `Decorator` class is made a subclass of `Component`. And since the `Decorator` class is unique in the design and cannot be

a subclass of a `ConcreteComponent` class, there is nothing to be gained by allowing intermediate classes in the hierarchy between the `Component` class and the `Decorator` class. Thus, we require that the `Component` class is a direct parent of the `Decorator` class, that is there is an inheritance relationship which links them. This property is expressed formally by the function *has\_parent\_direct* from [6], and the required property is then:

**value**

`has_parent_component` : `Wf_Design_Renaming`  $\rightarrow$  **Bool**  
`has_parent_component(dr)`  $\equiv$  `has_parent_direct(Decorator, Component, dr)`

The ability to recursively add decorators is achieved by storing in each decorator a reference to the object it decorates. This is represented by the state variable `component` and its associated aggregation relation linking the `Decorator` class to the `Component` class. Since decorators are added one at a time in a nested fashion rather than in a single group all at the same time, this variable is unique. In addition, the relation linking the `Decorator` class to the `Component` class is also unique since the role of the `Component` class is simply to provide a common, abstract interface to both the `Decorator/ConcreteDecorator` classes and the `ConcreteComponent` classes and the relation exists solely to implement the recursive composition of decorators.

Again these properties are analogous to the properties of the `adaptee` state variable in the `Adapter` class in the Object Adapter pattern and its corresponding relation with the `Adaptee` class (see Section 2.2.2), so they are similarly specified using the functions *store\_unique\_vble*, *has\_assoc\_aggr\_var\_ren* and *has\_unique\_assoc\_aggr\_relation*.

**value**

`store_unique_component` : `Wf_Design_Renaming`  $\rightarrow$  **Bool**  
`store_unique_component(dr)`  $\equiv$  `store_unique_vble(Decorator, component, dr)`,  
  
`decorator_relation`: `Wf_Design_Renaming`  $\rightarrow$  **Bool**  
`decorator_relation(dr)`  $\equiv$   
`has_assoc_aggr_var_ren(Decorator, Component, Aggregation, component, G.one, dr)  $\wedge$`   
`has_unique_assoc_aggr_relation(Decorator, Component, dr)`

According to the description of the participants of the pattern (see Section 5.1 above), the `Component` class defines the common interface for objects that can be decorated dynamically, that is for the `ConcreteComponent` classes. In addition, the `Decorator` class is described as having an interface which conforms to that of the `Component` class. This interface is represented by the methods which play the `Operation` role, and of course it may actually consist of many methods.

At the level of the `Component` class, all these methods are abstract since this class presents the common interface for all components and decorators and is the one through which clients interact. These properties are identical to the properties of the `Request` methods in the `Target` class in

the Adapter pattern (both versions) so are similarly specified using the functions *has\_def\_method* and *has\_all\_def\_method*.

**value**

```
Ct_has_operation_defined : Wf_Design_Renaming → Bool
Ct_has_operation_defined(dr) ≡
  has_def_method(Component, Operation, dr) ∧
  has_all_def_method(Component, Operation, dr)
```

The classes playing the **ConcreteComponent** role are responsible for implementing the operations defined in the **Component** class and thus for providing the actual behaviour the clients see. However, since the **ConcreteComponent** classes are subclasses of the **Component** class, the fact that the **Component** class contains at least one method playing the **Operation** role implies that there must also be at least one such method in each **ConcreteComponent** class. And the fact that the **ConcreteComponent** classes are concrete means that they cannot contain abstract methods, though they could contain error methods. It is therefore sufficient to specify that all **Operation** methods are implemented in the **ConcreteComponent** classes, which is done using the function *has\_all\_impl\_method*. This ensures that every component is able to respond to requests sent by clients. However, each design class playing the **ConcreteComponent** role represents a different component that can be decorated, so the actual implementations of these **Operation** methods is likely to be different in different classes. This means that we cannot specify anything about their bodies.

**value**

```
CCt_has_impl_operation : Wf_Design_Renaming → Bool
CCt_has_impl_operation(dr) ≡
  has_all_impl_method(ConcreteComponent, Operation, dr)
```

The **Operation** methods are also all implemented at the level of both the **Decorator** class and the **ConcreteDecorator** classes. In the case of the **Decorator** class, the methods simply forward the requests to the underlying component recursively by invoking the same method on the component state variable. This behaviour is analogous to that of the **Request** method in the **Adapter** class in the Object Adapter pattern (see Section 2.2.2) so is similarly specified by the function *deleg\_with\_var*.

**value**

```
Dec_has_impl_operation: Wf_Design_Renaming → Bool
Dec_has_impl_operation(dr) ≡
  deleg_with_var(Decorator, Operation, component, Component, Operation, dr)
```

The **ConcreteDecorator** classes also implement all the **Operation** methods, which is specified using the function *has\_all\_impl\_method* as for the **ConcreteComponent** classes:

**value**

```
CDec_has_impl_operation : Wf_Design_Renaming → Bool
CDec_has_impl_operation(dr) ≡
  has_all_impl_method(ConcreteDecorator, Operation, dr)
```

However, each of these classes may additionally tailor the methods to the particular decorator that it represents, adding functionality appropriate to that decorator to the basic recursive call mechanism implemented in the `Decorator` class. This may be done in two different ways.

One possibility is that the `ConcreteDecorator` introduces some new state variables, which are represented in the structure shown in Figure 9 by the `addedState` variable in the `ConcreteDecoratorA` class. In this case, the implementation of the `Operation` methods in the `ConcreteDecorator` class involves an invocation of the corresponding basic `Operation` method from the `Decorator` class, which implements the recursive element of the method, together with some additional functionality which depends on the `addedState` variables. In this way, the appropriate version of the `Operation` method is applied to each decorator in turn, then finally to the component which they decorate.

This type of `ConcreteDecorator` class must thus contain at least one state variable playing the `addedState` role and all its `Operation` methods must include a super invocation to the same `Operation`. The first of these properties is specified using the function *store\_vble* and the second by the function *exists\_super\_invocation*. Note however that we cannot specify the part of the behaviour of the `Operation` method that depends on the `addedState` variable because in general this will be different for different decorators. The specifications corresponding to this alternative are thus:

**value**

```
CDec_stores_added_state: Wf_Design_Renaming → Bool
CDec_stores_added_state(dr) ≡ store_vble(ConcreteDecorator, addedState, dr),

CDec_has_super_operation : Wf_Design_Renaming → Bool
CDec_has_super_operation(dr) ≡
  exists_super_invocation(ConcreteDecorator, Operation, dr)
```

The second way of extending the functionality of the `Operation` methods is to introduce some new methods into the `ConcreteDecorator` class. These are represented in the structure shown in Figure 9 by the `AddedBehaviour` method in the `ConcreteDecoratorB` class. In this case too, the implementation of the `Operation` methods in the `ConcreteDecorator` class involves an invocation of the corresponding basic `Operation` method from the `Decorator` class, which implements the recursive element of the method, but this time the additional functionality consists of invocations of the `AddedBehaviour` methods as shown in the annotation to the `Operation` method in Figure 9.

This type of `ConcreteDecorator` class must thus contain at least one method playing the `AddedBehaviour` role and all its `Operation` methods must include a super invocation to the same `Operation` as well as at least one (self) invocation to some method playing this `AddedBehaviour` role. Furthermore, all methods playing the `AddedBehaviour` role must be implemented, though we cannot specify anything about their implementation because again this will in general be different for different decorators.

The first and last of these properties are specified in the standard way using the functions *has\_impl\_method* and *has\_all\_impl\_method*, while the second is specified using *exists\_super\_self\_inv*. The specifications corresponding to this alternative are therefore:

**value**

```

CDec_has_impl_added_behaviour: Wf_Design_Renaming → Bool
CDec_has_impl_added_behaviour(dr) ≡
  has_impl_method(ConcreteDecorator, AddedBehaviour, dr) ∧
  has_all_impl_method(ConcreteDecorator, AddedBehaviour, dr),

CDec_has_super_self_operation : Wf_Design_Renaming → Bool
CDec_has_super_self_operation(dr) ≡
  exists_super_self_inv(ConcreteDecorator, Operation, AddedBehaviour, dr)

```

Any given `ConcreteDecorator` class must have a behaviour which conforms to at least one of the possibilities described above, and possibly but not necessarily both. These alternatives are expressed by the function *CDec\_has\_extended\_interface* – each class must add either new state variables or new methods and its `Operation` methods must invoke the basic `Operation` defined in the `Decorator` class through a super invocation; and if the class adds new methods at least one of these must also be invoked in each `Operation` method.

**value**

```

CDec_has_extended_interface : Wf_Design_Renaming → Bool
CDec_has_extended_interface(dr) ≡
  ( CDec_has_impl_added_behaviour(dr) ∨ CDec_stores_added_state(dr) ) ∧
  ( CDec_has_impl_added_behaviour(dr) ⇒ CDec_has_super_self_operation(dr) ) ∧
  ( CDec_stores_added_state(dr) ⇒ CDec_has_super_operation(dr) )

```

The structure of the Decorator pattern shown in Figure 9 does not include a `Client` class. However, the `Client` does actually play a specific role in the pattern, namely it interacts with the `Component` class by invoking its `Operation` interface. This means that the `Client` has a relation, which may be either an association or an aggregation, with the `Component` class. These properties are precisely analogous to the properties of the `Client` class in the Bridge pattern (see Section 3.2) so are similarly specified using the functions *has\_assoc\_aggr\_reltypes* and *use\_interface*.

**value**

```
decorator_client : Wf_Design_Renaming → Bool
decorator_client(dr) ≡
  has_assoc_aggr_retype(Client, Component, AssAggr, G.one, dr) ∧
  use_interface(Client, Component, Operation, dr)
```

Putting all these elements together we arrive at the following specification of the function *is\_decorator\_pattern* which checks whether a given design matches the Decorator pattern:

**value**

```
is_decorator_pattern : Wf_Design_Renaming → Bool
is_decorator_pattern(dr) ≡
  one_component_in_hierarchy(dr) ∧
  Ct_has_operation_defined(dr) ∧
  exist_concrete_component(dr) ∧
  is_concrete_component(dr) ∧
  CCt_has_impl_operation(dr) ∧
  one_decorator_in_hierarchy(dr) ∧
  is_abstract_decorator(dr) ∧
  has_parent_component(dr) ∧
  store_unique_component(dr) ∧
  Dec_has_impl_operation(dr) ∧
  decorator_client(dr) ∧
  decorator_relation(dr) ∧
  exist_concrete_decorator(dr) ∧
  is_concrete_decorator(dr) ∧
  CDec_has_impl_operation(dr) ∧ CDec_has_extended_interface(dr)
```

## 6 The Facade Pattern

Structuring a system into subsystems helps reduce complexity, but it is also useful to minimise the communication and dependencies between subsystems. One way to achieve this is to introduce an object that provides a single, simplified interface to the more general facilities of a subsystem. This object thus acts as a “facade” for the subsystem, providing a higher-level interface that can shield clients from the complexities of the subsystem classes themselves. The facade thus offers clients some specifically tailored but possibly restricted implementation(s) of the functionality of the subsystems [7]. The Facade pattern describes a way in which this can be achieved.

As usual, we first present the essential properties of the pattern as described in [7], then go on to give our analysis and formal specification.

## 6.1 Properties of the Facade Pattern

### Intent

Provide a unified interface to a set of interfaces in a subsystem. Facade defines a higher-level interface that makes the subsystem easier to use.

### Structure

Figure 11 shows the structure of the Facade pattern.

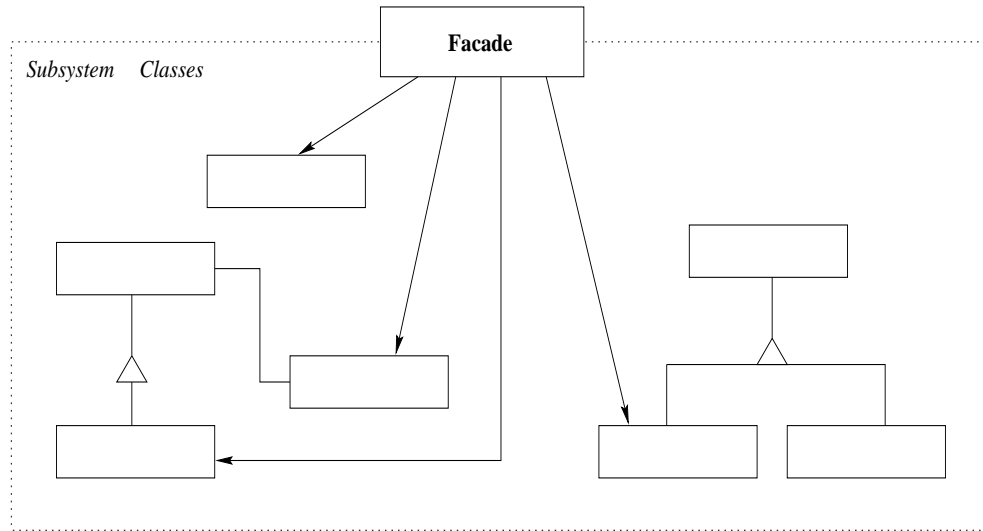


Figure 11: Facade Pattern Structure

### Participants

- Facade
  - knows which subsystem classes are responsible for a request.
  - delegates client requests to appropriate subsystem objects.
- subsystem classes
  - implement subsystem functionality.
  - handle work assigned by the Facade object.
  - have no knowledge of the facade; that is, they keep no references to it.

### Collaborations

- Clients communicate with the subsystem by sending requests to Facade, which forwards them to the appropriate subsystem object(s). Although the subsystem objects perform the actual work, the facade may have to do work of its own to translate its interface to subsystem interfaces.
- Clients that use the facade don't have to access its subsystem objects directly.



## 6.2 Formalising the Facade Pattern

The Facade pattern structure shown in Figure 11 does not give a clear representation of the pattern; that is, it does not show properly the behaviour of the participant classes, in particular the subsystem classes. However, some essential properties can be elicited from the description of the pattern given above.

First of all, the description of the participants of the pattern clearly indicates that subsystem classes play an important role in the pattern, so we introduce an explicit class role **SubsystemClass** to represent this. We also introduce a method role called **HandleRequest** into the interface of this **SubsystemClass** role to correspond to its responsibility for handling work assigned to it by the Facade object.

In order to make this delegation of responsibility, the **Facade** class must hold references to the subsystems it represents. We introduce the state variable **subsystemClass** to represent these references, and a corresponding association relation linking the **Facade** class to the various instances of **SubsystemClass**. We further introduce a new method role **Request** into the **Facade** interface to represent the operations which perform the delegation, this being done by simply invoking the appropriate **HandleRequest** operation(s) on the appropriate **subsystemClass** state variable(s).

The **Facade** class may also be responsible for creating instances of the **SubsystemClass** classes to satisfy demands from clients, so we introduce another method role **Creator** into the **Facade** class to model this instantiation. The result of this method should be an instance of some **SubsystemClass**.

The role of clients in the pattern is also important, as explained in the description of the pattern's collaborations, so we introduce another new class role **Client** to represent the client. We also introduce an association relation linking the **Client** to the **Facade** class to model the fact that clients using the **Facade** do not need to access subsystem classes directly but can rather interact with them solely through the **Facade** interface.

One of the main goals of the pattern is to minimise dependencies in the way clients communicate with subsystems, and the discussion of the implementation of the pattern in [7] suggests another way of reducing the coupling between the **Client** and the subsystems which is applicable when the different subsystems can be combined in different ways to provide clients with different functionality. In such a situation, it can be useful to make the **Facade** class an abstract class with concrete subclasses for each of the different implementations (i.e. combinations of subsystems) offered to the **Client**. To model this, we introduce one more new class role which we call **SubclassFacade** and which represents concrete subclasses of the **Facade** class. The **SubclassFacade** classes, if they exist, should of course implement the interface defined in the **Facade** class, which in this case will be abstract.

These considerations lead us to the more detailed structure of the Facade pattern shown in Figure 12, and it is this extended structure on which we base our analysis and formal specification

below.

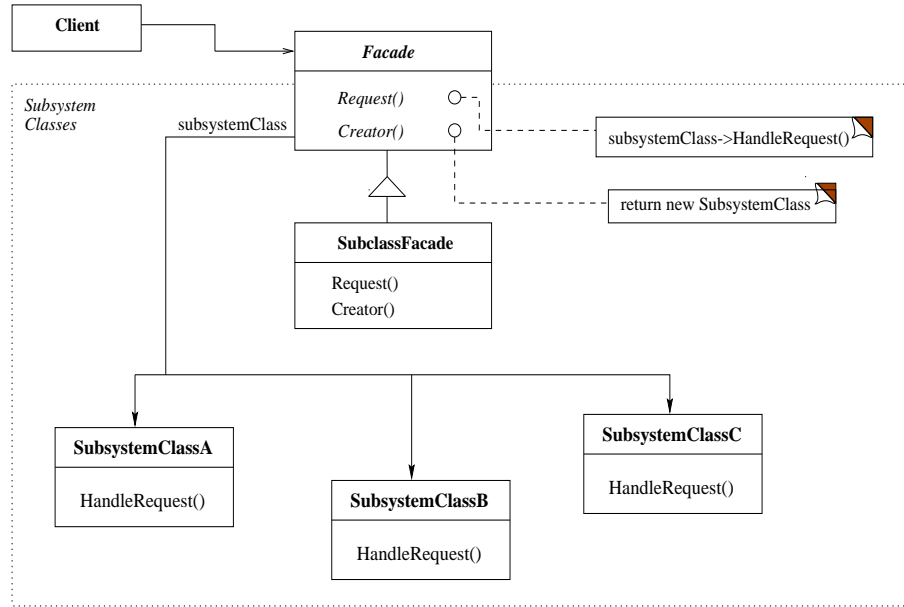


Figure 12: The Modified Facade Pattern Structure

We begin in the standard way by introducing RSL constants representing all the entities in the (modified) structure.

**value**

Facade, SubclassFacade, SubsystemClass, Client : G.Class\_Name,

Creator, Request, HandleRequest : G.Method\_Name,

subsystemClass : G.Variable\_Name

The classes in the hierarchy of facade classes can only play the roles **Facade** or **SubclassFacade**, and we can additionally assume without loss of generality that there is only a single such hierarchy, and hence only a single **Facade** class in the design, as usual considering that a design with more than one **Facade** class corresponds to more than one instance of the pattern. This property is once more specified using the function *hierarchy*.

**value**

one\_facade\_in\_hierarchy : Wf\_Design\_Renaming  $\rightarrow$  **Bool**

one\_facade\_in\_hierarchy(dr)  $\equiv$

hierarchy(Facade, {SubclassFacade}, {SubsystemClass, Client}, dr)

In this case, however, the **Facade** class may be abstract or concrete: it is abstract if there are different possible combinations of the subsystem classes which can be offered to the client, in which case each such combination is represented by a particular class playing the **SubclassFacade** role; on the other hand it is concrete if there is only a single possible combination of subsystem classes, in which case the role **SubclassFacade** is redundant and does not appear in the design.

We therefore simply specify that the **Facade** class must be concrete if there is no class playing the **SubclassFacade** role:

```

value
  F_concrete_or_SF_exists : Wf_Design_Renaming → Bool
  F_concrete_or_SF_exists(dr) ≡
    is_concrete_class(Facade, dr) ∨ exists_role(SubclassFacade, dr)

```

Of course if the design does contain classes playing the **SubclassFacade** role all these classes should be concrete subclasses of the class playing the **Facade** role. As for the other patterns, this property is specified using the function *is\_concrete*.

```

value
  SF_is_a_concrete_F : Wf_Design_Renaming → Bool
  SF_is_a_concrete_F(dr) ≡ is_concrete(Facade, SubclassFacade, dr)

```

Since the **Facade** class effectively “hides” the subsystem classes from clients and must forward requests from clients to the appropriate subsystem object(s), it must contain a reference to at least one such object, that is it must have at least one state variable playing the **subsystemClass** role. In addition, this variable must correspond to a relation with a class playing the **SubsystemClass** role, though this could be an association or an aggregation relation depending on context. However, neither the state variable nor the relation is necessarily unique since the **Facade** class may need to reference more than one **SubsystemClass** or even more than one instance of the same **SubsystemClass**. The above properties are therefore specified simply using the functions *store\_vble* and *has\_assoc\_aggr\_var\_ren*.

```

value
  F_stores_subsystemclasses : Wf_Design_Renaming → Bool
  F_stores_subsystemclasses(dr) ≡ store_vble(Facade, subsystemClass, dr),

  F_knows_all_SC : Wf_Design_Renaming → Bool
  F_knows_all_SC(dr) ≡
    has_assoc_aggr_var_ren(Facade, SubsystemClass, AssAggr, subsystemClass, G.one, dr)

```

There may be many methods playing the **Request** role in the **Facade** and **SubclassFacade** classes, though there should be at least one such method since this represents the interface used by clients. If the design includes classes playing the **SubclassFacade** role, the methods playing the **Request** role must be implemented in these classes but they may be either abstract or concrete in the **Facade** class. If there are no classes playing the **SubclassFacade** role in the design the **Request** methods must be implemented in the **Facade** class. Either way, the body of each implemented **Request** method must include an invocation to a variable playing the **subsystemClass** role of a method playing the **HandleRequest** role in a class playing the **SubsystemClass** role.

The function *deleg\_var\_to\_some\_class* from [6] states that if a method playing a given role (**Request**) in a given class (**Facade**) is implemented, then its body includes an invocation to a variable playing a given role (**subsystemClass**) of a method playing a second given role (**HandleRequest**). The function *exists\_method* specifies that the **Facade** class contains at least one method playing the **Request** role, and the function *has\_error\_method* is used in the negative to ensure that all **Request** methods are either defined or implemented. In the case where there is no class playing the **SubclassFacade** role and the **Facade** class is concrete, this then implies that the **Request** methods are all implemented in the **Facade** class and that the body of each contains the required invocation of a **HandleRequest** method.

**value**

```
F_includes_request: Wf_Design_Renaming → Bool
F_includes_request(d, r) ≡
  exists_method(Facade, Request, r) ∧
  ~has_error_method(Facade, Request, (d, r)) ∧
  deleg_var_to_some_class
    (Facade, Request, subsystemClass, SubsystemClass, HandleRequest, (d, r))
```

The specification of the **Request** methods in the **SubclassFacade** classes, if they exist, is similar except that here some but not all of the methods can be error methods since not all subclasses necessarily implement the same behaviour. We therefore use the function *has\_impl\_method* to specify that at least one **Request** method is implemented, and the function *deleg\_var\_to\_some\_class* is again used to ensure that all implemented **Request** methods perform the correct invocation as for the methods in the **Facade** class.

**value**

```
SF_has_impl_request : Wf_Design_Renaming → Bool
SF_has_impl_request(dr) ≡
  has_impl_method(SubclassFacade, Request, dr) ∧
  deleg_var_to_some_class
    (SubclassFacade, Request, subsystemClass, SubsystemClass, HandleRequest, dr)
```

As mentioned above, the **Facade** class may be responsible for creating instances of the **SubsystemClass** classes to satisfy demands from clients. This possibility is represented by the method

role **Creator** which returns an instance of some **SubsystemClass** as its result.

The specification of the properties of this method in both the **Facade** class and the **SubclassFacade** subclasses is similar to that of the **Request** methods discussed above. Again, there may be many methods playing the **Creator** role in the **Facade** and **SubclassFacade** classes, though there should be at least one such method, and if the design includes classes playing the **SubclassFacade** role, the methods playing the **Creator** role must be implemented in these classes but they may be either abstract or concrete in the **Facade** class. Furthermore, if there are no classes playing the **SubclassFacade** role in the design the **Creator** methods must be implemented in the **Facade** class, and in any case each implemented **Creator** method must return a result which is some instance of a class playing the **SubsystemClass** role.

In the formal specification the function *has\_method\_with\_result\_class* from [6] forms the analogue of the function *deleg\_var\_to\_some\_class* which was used above in the specification of the properties of the **Request** methods. It states that if a method playing a given role (**Creator**) in a given class (**Facade**) is implemented, then its result must be a variable representing a class playing a second given role (**SubsystemClass**). And the functions *exists\_method* and *has\_error\_method* are again used to specify that the **Facade** class contains at least one method playing the **Creator** role and all methods playing this role are either defined or implemented in this class. Once again, in the case where there is no class playing the **SubclassFacade** role and the **Facade** class is concrete, this then implies that the **Creator** methods are all implemented in the **Facade** class and that their result is an appropriate instance of some class playing the **SubsystemClass** role. Note however that we additionally use the function *has\_method\_with\_res* to ensure that the method always returns some result even if this is not determined precisely when the **Facade** class is abstract.

In the case of classes playing the **SubclassFacade** role, if they exist, once again some but not all of the **Creator** methods can be error methods, again because not all subclasses necessarily implement the same behaviour. So we analogously use the function *has\_impl\_method* to specify that at least one **Creator** method is implemented in each such class, then the function *has\_method\_with\_result\_class* similarly ensures that all implemented **Creator** methods have the correct result as in the case of the **Facade** class. Note that in this case we do not need to include the function *has\_method\_with\_res* since this was already specified for the **Facade** class and the general model in [6] requires that the result of an operation is consistent within a hierarchy of classes.

**value**

$F\_includes\_creator : Wf\_Design\_Renaming \rightarrow \mathbf{Bool}$

$F\_includes\_creator(d, r) \equiv$

$exists\_method(Facade, Creator, r) \wedge$

$\sim has\_error\_method(Facade, Creator, (d, r)) \wedge$

$has\_method\_with\_res(Facade, Creator, (d, r)) \wedge$

$has\_method\_with\_result\_class(Facade, Creator, SubsystemClass, (d, r)),$

$SF\_has\_impl\_creator : Wf\_Design\_Renaming \rightarrow \mathbf{Bool}$

$$\begin{aligned} \text{SF\_has\_impl\_creator}(\text{dr}) \equiv \\ & \text{has\_impl\_method}(\text{SubclassFacade}, \text{Creator}, \text{dr}) \wedge \\ & \text{has\_method\_with\_result\_class}(\text{SubclassFacade}, \text{Creator}, \text{SubsystemClass}, \text{dr}) \end{aligned}$$

In principle every **SubsystemClass** may be instantiated by the **Facade** class using a method playing the **Creator** role since the **Facade** class knows about all the subsystem classes and may create objects from these classes when appropriate. This means that there must be an instantiation relation connecting the **Facade** class to each **SubsystemClass**. This is specified using the function *has\_instantiation*.

However the converse is not true, and in fact it is explicitly stated in the description of the pattern's participants above that the subsystem classes have no knowledge of the **Facade** class or of its subclasses. In particular they keep no references to any facade class, nor may they instantiate facade classes. Thus, there can be no relations linking a class playing the **SubsystemClass** role to a class playing either the **Facade** or the **SubclassFacade** role. This is specified using the function *not\_related\_classes*.

**value**

$$\begin{aligned} \text{SC\_instantiated\_by\_F} : \text{Wf\_Design\_Renaming} &\rightarrow \mathbf{Bool} \\ \text{SC\_instantiated\_by\_F}(\text{dr}) &\equiv \text{has\_instantiation}(\text{Facade}, \text{SubsystemClass}, \text{dr}), \end{aligned}$$

$$\begin{aligned} \text{SC\_dont\_know\_about\_F} : \text{Wf\_Design\_Renaming} &\rightarrow \mathbf{Bool} \\ \text{SC\_dont\_know\_about\_F}(\text{dr}) &\equiv \\ & \text{not\_related\_classes}(\text{SubsystemClass}, \text{Facade}, \text{dr}) \wedge \\ & \text{not\_related\_classes}(\text{SubsystemClass}, \text{SubclassFacade}, \text{dr}) \end{aligned}$$

The **SubsystemClass** role represents the classes that handle requests forwarded by the **Facade** class so in general there should be at least one such class and all such classes should be concrete. In addition, each should include at least one implemented method playing the **HandleRequest** role. These properties are specified using the functions *exists\_role*, *is\_concrete\_class* and *has\_impl\_method*.

**value**

$$\begin{aligned} \text{exists\_subclass} : \text{Wf\_Design\_Renaming} &\rightarrow \mathbf{Bool} \\ \text{exists\_subclass}(\text{dr}) &\equiv \text{exists\_role}(\text{SubsystemClass}, \text{dr}), \\ \\ \text{subclass\_concrete} : \text{Wf\_Design\_Renaming} &\rightarrow \mathbf{Bool} \\ \text{subclass\_concrete}(\text{dr}) &\equiv \text{is\_concrete\_class}(\text{SubsystemClass}, \text{dr}), \\ \\ \text{SC\_has\_handleRequest} : \text{Wf\_Design\_Renaming} &\rightarrow \mathbf{Bool} \\ \text{SC\_has\_handleRequest}(\text{dr}) &\equiv \text{has\_impl\_method}(\text{SubsystemClass}, \text{HandleRequest}, \text{dr}) \end{aligned}$$

Clients communicate with subsystems through the interface of the Facade class as represented by the relation linking these two classes in the modified structure shown in Figure 12. This reference is usually an association because in many cases both the facade and the subsystems have life cycles independent of clients, for example when there is only one instance of the Facade class serving all possible clients, but an aggregation is possible in some cases. In addition, the Client class communicates with the Facade class by invoking Request methods in its interface.

These properties are exactly analogous to the properties of the Client role in the Bridge pattern (see Section 3.2) so they are also formally specified using the functions *has\_assoc\_aggr\_reltype* and *use\_interface*.

**value**

```

facade_client : Wf_Design_Renaming → Bool
facade_client(dr) ≡
  has_assoc_aggr_reltype(Client, Facade, AssAggr, G.one, dr) ∧
  use_interface(Client, Facade, Request, dr)

```

The conjunction of all the above properties then yields the function *is\_facade* which checks whether a given design matches the Facade pattern.

**value**

```

is_facade : Wf_Design_Renaming → Bool
is_facade(dr) ≡
  one_facade_in_hierarchy(dr) ∧
  exists_subsystemclass(dr) ∧
  subsystemclass_concrete(dr) ∧
  F_concrete_or_SF_exists(dr) ∧
  facade_client(dr) ∧
  SF_is_a_concrete_F(dr) ∧
  F_includes_request(dr) ∧
  SF_has_impl_request(dr) ∧
  SC_has_handleRequest(dr) ∧
  F_stores_subsystemclasses(dr) ∧
  F_knows_all_SC(dr) ∧
  SC_dont_know_about_F(dr)

```

## 7 The Flyweight Pattern

When saving space is important, it is helpful to reduce the number of objects an application requires as much as possible.

Some design structures naively require a big object structure at run time. This can often be reduced by sharing objects where possible, but although this is a good approach it is not always enough. In such cases it is useful to look at objects from other point of view: since objects of the same class share the same behaviour, the difference between them is related only to their properties and separating the properties into those which are common to all objects in the class and those which depend on the context in which they are applied may offer better opportunities for sharing.

A *flyweight* is a shared object that can be used in multiple contexts simultaneously. It acts as an independent object in each context – it is indistinguishable from an instance of the object that is not shared – and it cannot make assumptions about the context in which it operates. The important concept here is a distinction between intrinsic and extrinsic state. Intrinsic state is stored in the flyweight and consists of information that is independent of the flyweight’s context, thereby making the flyweight sharable. Extrinsic state, on the other hand, depends on and varies with the flyweight’s context and therefore cannot be shared. Client objects are responsible for passing extrinsic state to the flyweight when it needs it [7].

The Flyweight pattern describes a design in which objects are shared using flyweights. We first introduce the essential elements of the pattern as described in [7], then we present our analysis and formal specification of their properties.

## 7.1 Properties of the Flyweight Pattern

### Intent

Use sharing to support large numbers of fine-grained objects efficiently.

### Structure

Figures 13 and 14 show respectively the structure of the Flyweight pattern and an object diagram illustrating how flyweights can be shared.

### Participants

- Flyweight
  - declares an interface through which flyweights can receive and act on extrinsic state.
- ConcreteFlyweight
  - implements the Flyweight interface and adds storage for intrinsic state, if any. ConcreteFlyweight object must be sharable. Any state it stores must be intrinsic; that is, it must be independent of the ConcreteFlyweight object’s context.
- UnsharedConcreteFlyweight
  - not all Flyweight subclasses need to be shared. The Flyweight interface enables sharing; it doesn’t enforce it. It’s common for UnsharedConcreteFlyweight ob-



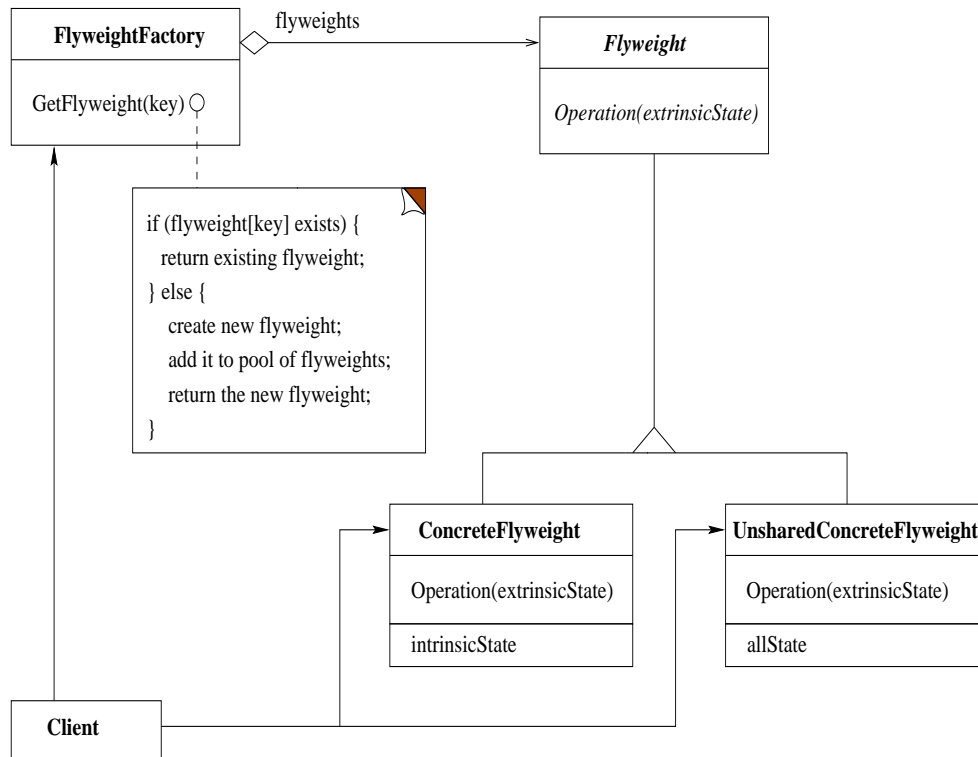


Figure 13: Flyweight Pattern Structure

jects to have ConcreteFlyweight objects as children at some level in the flyweight object structure.

- FlyweightFactory
  - creates and manages flyweight objects.
  - ensures that flyweights are shared properly. When a client requests a flyweight, the FlyweightFactory object supplies an existing instance or creates one, if none exists.
- Client
  - maintains a reference to flyweight(s).
  - computes or stores the extrinsic state of flyweight(s).

### Collaborations

- State that a flyweight needs to function must be characterised as either intrinsic or extrinsic. Intrinsic state is stored in the ConcreteFlyweight object; extrinsic state is stored or computed by Client objects. Clients pass this state to the flyweight when they invoke its operations.
- Clients should not instantiate ConcreteFlyweights directly. Clients must obtain ConcreteFlyweight objects exclusively from the FlyweightFactory object to ensure they are shared properly.

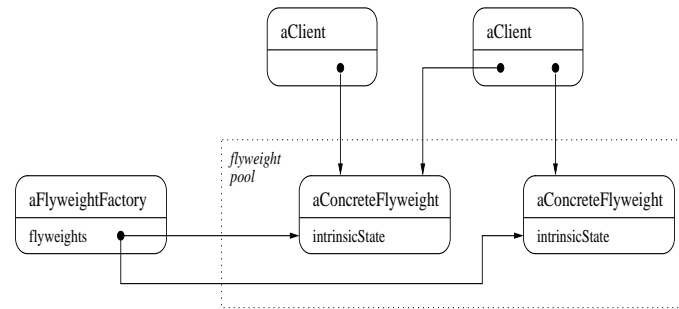


Figure 14: Flyweight Object Structure

## 7.2 Formalising the Flyweight Pattern

Although no details regarding the functionality of the Client class are included in the structure of the pattern shown in Figure 13, its responsibilities are defined in some detail in the description of the participants and collaborations of the Flyweight pattern above: it must pass extrinsic state to the flyweights when it invokes their operations, and it must interact correctly with the FlyweightFactory so as to ensure that flyweights are shared correctly. We therefore begin by extending the structure as shown in Figure 15 to explicitly include these interactions.

As stated in the collaborations of the pattern, the Client may only obtain a ConcreteFlyweight from the FlyweightFactory. In fact it does this by invoking the GetFlyweight method in the FlyweightFactory class, the purpose of which is to ensure that the sharing of flyweights is handled correctly: if a client asks for a flyweight which already exists this existing flyweight is returned; otherwise a new flyweight is created and recorded, then this new flyweight is returned.

However, this procedure applies to shared but not to unshared flyweights: multiple instances of the same unshared flyweight may exist; indeed in the case of unshared flyweights the Client would normally create a new object every time since unshared flyweights are not intended to be reused. We therefore introduce the new method role CreateFlyweight into the FlyweightFactory class to represent the method the Client invokes to create unshared flyweights. This method is therefore the analogue of the GetFlyweight method for unshared flyweights and it simply returns a new instance of the appropriate UnsharedConcreteFlyweight class as shown by its annotation.

We also introduce two methods, SMethod and UMethod, into the Client class to describe its interaction with shared and unshared flyweights respectively. As shown in the annotations to these methods, they first invoke the appropriate “creation” method (GetFlyweight or CreateFlyweight respectively) to obtain the flyweight, then the Client can invoke operations on these flyweights, passing any necessary extrinsic state as parameters to the invocations as stated in the collaborations above. These interactions are shown in the annotations to the methods SMethod and UMethod in Figure 15.

We use this extended structure as the basis for our analysis and specification. We begin as usual by defining RSL constants representing the entities in this structure:

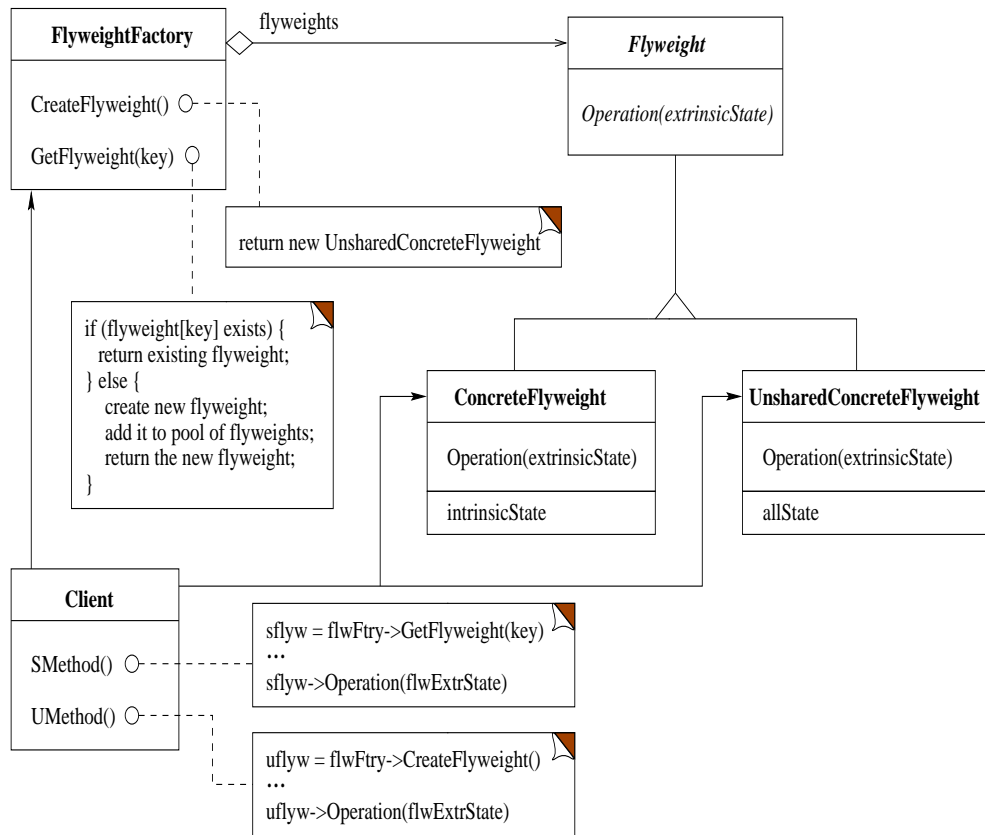


Figure 15: Flyweight Pattern Structure

**value**

Flyweight, FlyweightFactory, ConcreteFlyweight,  
UnsharedConcreteFlyweight, Client : G.Class\_Name,

key, flyweights, intrinsicState, extrinsicState, allState : G.Variable\_Name,

GetFlyweight, CreateFlyweight, Operation, SMethod, UMethod : G.Method\_Name

The structure of the pattern consists essentially of the hierarchy of flyweight classes (Flyweight, ConcreteFlyweight and UnsharedConcreteFlyweight) together with the FlyweightFactory, which keeps a record of existing (shared) concrete flyweights, and the Client, which invokes the CreateFlyweight and GetFlyweight methods in the FlyweightFactory class and the Operation methods in the concrete flyweight classes.

In the hierarchy of flyweight classes, the behaviour and properties common to both kinds of flyweight (shared and unshared) is factored out in the Flyweight class, and leaf classes represent concrete flyweights and therefore play either the ConcreteFlyweight or the UnsharedConcreteFlyweight role. In addition, a flyweight cannot be both shared and unshared, so no class in the

design can play both the `ConcreteFlyweight` and the `UnsharedConcreteFlyweight` role, nor is it possible for a `ConcreteFlyweight` class to be either a superclass or a subclass of an `UnsharedConcreteFlyweight` class. In addition, no class in the hierarchy can play either the `Client` role or the `FlyweightFactory` role.

In addition, we can assume without loss of generality that each instance of the pattern has only one such hierarchy in the design, once again considering each such hierarchy as a separate pattern instance. This means that the `Flyweight` class can also be considered to be unique. It is also abstract since its only purpose is to define a common interface which allows clients to treat its subclasses uniformly, though we again do not need to specify this explicitly because the class contains an abstract method (`Operation`) and so must be abstract itself.

The above properties are specified as usual using the *hierarchy* function.

**value**

```
one_flyweight_in_hierarchy : Wf_Design_Renaming → Bool
one_flyweight_in_hierarchy(dr) ≡
  hierarchy
    (Flyweight, {ConcreteFlyweight, UnsharedConcreteFlyweight},
     {FlyweightFactory, Client}, dr)
```

For the classes representing the concrete flyweights, there can be many classes playing the `ConcreteFlyweight` role and also many classes playing the `UnsharedConcreteFlyweight` role. In addition, there must be at least one class playing the `ConcreteFlyweight` role, that is at least one shared flyweight, since the whole point of the Flyweight pattern is to describe the sharing of flyweight objects. However, although some flyweights may be unshared, it is in fact possible that they are all shared, so there may be no class in the design playing the `UnsharedConcreteFlyweight` role.

We specify the existence of the `ConcreteFlyweight` role in the usual way using the function *exists\_role*, but there is no corresponding specification for the `UnsharedConcreteFlyweight` role since this may or may not exist. We also specify that all the `ConcreteFlyweight` and `UnsharedConcreteFlyweight` classes should be concrete subclasses of the class playing the `Flyweight` role, again using the function *is\_concrete* as usual. The specification of the above properties is thus:

**value**

```
exists_concrete_flyweight : Wf_Design_Renaming → Bool
exists_concrete_flyweight(dr) ≡ exists_role(ConcreteFlyweight, dr),

are_concrete_flyweight : Wf_Design_Renaming → Bool
are_concrete_flyweight(dr) ≡
  is_concrete(Flyweight, ConcreteFlyweight, dr) ∧
  is_concrete(Flyweight, UnsharedConcreteFlyweight, dr)
```

According to the description of the participants of the pattern (see Section 7.1), the `Flyweight` class defines an interface through which flyweights can receive and act on extrinsic state. This interface is represented by the `Operation` methods, and the parameter `extrinsicState` of these methods represents the extrinsic state they receive.

The interface can of course consist of more than one `Operation` method, though there must be at least one such method otherwise there is no interface and hence no way of interacting with the flyweights.

At the level of the `Flyweight` class, the interface is only defined (abstract), so we use the functions *has\_def\_method* and *has\_all\_def\_method* to specify it as in the specification of the properties of the `Request` method in the `Target` class in both versions of the Adapter pattern. However, in this case we also need to specify that the `Operation` method has at least one parameter that plays the `extrinsicState` role, for which we use the function *fparams\_var\_ren* from [6]. The specification of the properties of the `Operation` methods in the `Flyweight` class is therefore:

**value**

```
F_has_defined_operation : Wf_Design_Renaming → Bool
F_has_defined_operation(dr) ≡
  has_def_method(Flyweight, Operation, dr) ∧
  has_all_def_method(Flyweight, Operation, dr) ∧
  fparams_var_ren(Flyweight, Operation, extrinsicState, dr)
```

Every subclass of `Flyweight` must implement the `Operation` interface, that is all such methods must be implemented in both the `ConcreteFlyweight` and the `UnsharedConcreteFlyweight` classes. These properties are again specified using the function *has\_all\_impl\_method*.

**value**

```
CF_has_impl_operation : Wf_Design_Renaming → Bool
CF_has_impl_operation(dr) ≡
  has_all_impl_method(ConcreteFlyweight, Operation, dr),

UCF_has_impl_operation : Wf_Design_Renaming → Bool
UCF_has_impl_operation(dr) ≡
  has_all_impl_method(UnsharedConcreteFlyweight, Operation, dr)
```

In addition, the concrete flyweight classes should store the intrinsic state of the flyweights. This is represented by the state variables `intrinsicState` and `allState` for the `ConcreteFlyweight` and `UnsharedConcreteFlyweight` classes respectively. In both cases the intrinsic state may actually be stored in more than one variable so we simply use the function *store\_vble* to specify that at least one such variable must exist in both cases.

**value**

`CF_stores_intrinsic_state : Wf_Design_Renaming → Bool`

`CF_stores_intrinsic_state(dr) ≡ store_vble(ConcreteFlyweight, intrinsicState, dr),`

`UCF_stores_all_state : Wf_Design_Renaming → Bool`

`UCF_stores_all_state(dr) ≡ store_vble(UnsharedConcreteFlyweight, allState, dr)`

Turning now to the `FlyweightFactory` role, this is responsible for the creation and management of all flyweight objects, that is objects of the `ConcreteFlyweight` and `UnsharedConcreteFlyweight` classes. Since we only have a single hierarchy of flyweights, we can similarly assume without loss of generality that there is also only one class in the design playing the `FlyweightFactory` role: again we consider a design containing more than one such class as comprising more than one instance of the Flyweight pattern. This property is expressed formally using the function *exists\_one*.

**value**

`exists_one_flyweight_factory : Wf_Design_Renaming → Bool`

`exists_one_flyweight_factory(dr) ≡ exists_one(FlyweightFactory, dr)`

The `FlyweightFactory` class is also responsible for recording the existing shared flyweights so that it can ensure that clients receive the existing ones instead of new ones when they request shared flyweights. These existing flyweights are stored in the state variable `flyweights`, which has an associated aggregation relation linking the `FlyweightFactory` class to the `Flyweight` class. One such variable, and hence also one aggregation relation, is sufficient because all shared flyweights from all `ConcreteFlyweight` classes could be stored in a single collection. However, it is also possible that the shared flyweights are stored in more than one variable, for instance there might be one such variable for each separate `ConcreteFlyweight` class, in which case there would also be more than one aggregation relation. We therefore simply specify that there must be at least one state variable playing the `flyweights` role and at least one aggregation relation with the same name. These properties are specified respectively using the functions *store\_vble* and *has\_assoc\_aggr\_var\_ren*.

**value**

`FF_aggregates_flyweights : Wf_Design_Renaming → Bool`

`FF_aggregates_flyweights(dr) ≡`

`has_assoc_aggr_var_ren(`

`FlyweightFactory, Flyweight, Aggregation, flyweights, G.many, dr),`

`FF_stores_flyweights : Wf_Design_Renaming → Bool`

`FF_stores_flyweights(dr) ≡ store_vble(FlyweightFactory, flyweights, dr)`

The methods `CreateFlyweight` and `GetFlyweight` represent the behaviour of the `FlyweightFactory` class when the client requests an unshared or a shared flyweight respectively.

The `CreateFlyweight` method simply instantiates the appropriate `UnsharedConcreteFlyweight` class and returns the new instance thus created. However, since a design does not necessarily include unshared flyweights, the `FlyweightFactory` class does not necessarily include this method. Thus, we simply use the function *has\_all\_impl\_method* to specify that all methods playing this role must be implemented, and the function *has\_method\_with\_result\_class* to specify that the result of each such method is a new instance of an `UnsharedConcreteFlyweight` class.

The `GetFlyweight` method, on the other hand, must exist since there must be at least one `ConcreteFlyweight` class. This method first checks whether the requested flyweight, which is identified by the parameter `key`, already exists (i.e. belongs to the collection of flyweights stored in one of the `flyweights` state variables), and if so it returns this existing flyweight. Otherwise it creates a new flyweight by instantiating the appropriate `ConcreteFlyweight` class, adds this to the collection of existing flyweights and returns it. The existence of the `GetFlyweight` method is specified using the function *has\_impl\_method* as usual, and the function *fparams\_var\_ren* specifies that the method has a parameter playing the `key` role. The properties of the body of the method are captured by the function *res\_chng\_vble\_alternative*.

In addition, we specify that the `FlyweightFactory` class has instantiation relations with both the `ConcreteFlyweight` and the `UnsharedConcreteFlyweight` classes, though these relations are not shown explicitly in the pattern structure. This is done using the function *has\_instantiation*.

**value**

`FF_has_impl_get_flyweight : Wf_Design_Renaming → Bool`

`FF_has_impl_get_flyweight(dr) ≡`  
`has_impl_method(FlyweightFactory, GetFlyweight, dr) ∧`  
`fparams_var_ren(FlyweightFactory, GetFlyweight, key, dr) ∧`  
`res_chng_vble_alternative`  
`(FlyweightFactory, GetFlyweight, flyweights, key,`  
`G.collectionelement, ConcreteFlyweight, G.collectionadd ,dr),`

`FF_has_impl_create_flyweight : Wf_Design_Renaming → Bool`

`FF_has_impl_create_flyweight(dr) ≡`  
`has_all_impl_method(FlyweightFactory, CreateFlyweight, dr) ∧`  
`has_method_with_result_class`  
`(FlyweightFactory, CreateFlyweight, UnsharedConcreteFlyweight, dr),`

`FF_creates_flyweights : Wf_Design_Renaming → Bool`

`FF_creates_flyweights(dr) ≡`  
`has_instantiation(FlyweightFactory, ConcreteFlyweight, dr) ∧`  
`has_instantiation(FlyweightFactory, UnsharedConcreteFlyweight, dr)`

The Client is related to the `ConcreteFlyweight`, the `UnsharedConcreteFlyweight` and the `FlyweightFactory` classes by association relations. Clients may not instantiate the `ConcreteFlyweight` classes directly since they must always obtain shared flyweights by invoking the `GetFlyweight` method in the `FlyweightFactory`. However, it is possible for them to instantiate `UnsharedConcreteFlyweight` classes directly since these objects are not shared and there are therefore no restrictions on their creation. We use the function *has\_assoc\_aggr\_retype* to specify the existence of the three association relations, and the function *has\_no\_instantiation* to ensure that the Client class has no instantiation relation with the `ConcreteFlyweight` classes.

**value**

```
flyweight_client : Wf_Design_Renaming → Bool
flyweight_client(dr) ≡
  has_assoc_aggr_retype(Client, ConcreteFlyweight, Association, G.one, dr) ∧
  has_no_instantiation(Client, ConcreteFlyweight, dr) ∧
  has_assoc_aggr_retype(Client, UnsharedConcreteFlyweight, Association, G.one, dr) ∧
  has_assoc_aggr_retype(Client, FlyweightFactory, Association, G.one, dr)
```

The Client obtains `ConcreteFlyweight` objects from the `FlyweightFactory` by invoking the `GetFlyweight` operation, then having obtained the appropriate flyweight in this way it can interact with that flyweight through its `Operation` interface, passing the required extrinsic state as parameters to the invocations. This functionality is represented by the `SMethod` method. The `UMethod` method, which applies to unshared flyweights, is similar to the `SMethod` method except that it invokes the `CreateFlyweight` method instead of the `GetFlyweight` method. The functionality of both these methods is expressed using the function *lvar\_chng\_inv\_deleg*.

**value**

```
client_invokes : Wf_Design_Renaming → Bool
client_invokes(dr) ≡
  lvar_chng_inv_deleg
    (Client, SMethod, FlyweightFactory,
     GetFlyweight, ConcreteFlyweight, Operation, dr) ∧
  lvar_chng_inv_deleg
    (Client, UMethod, FlyweightFactory,
     CreateFlyweight, UnsharedConcreteFlyweight, Operation, dr)
```

A design then matches the Flyweight pattern if it satisfies all the properties above. The function *is\_flyweight\_pattern* specifies this.

**value**

```
is_flyweight_pattern : Wf_Design_Renaming → Bool
is_flyweight_pattern(dr) ≡
```



$$\begin{aligned}
& \text{one\_flyweight\_in\_hierarchy}(\text{dr}) \wedge \\
& \text{exists\_one\_flyweight\_factory}(\text{dr}) \wedge \\
& \text{exists\_concrete\_flyweight}(\text{dr}) \wedge \\
& \text{flyweight\_client}(\text{dr}) \wedge \\
& \text{client\_invokes}(\text{dr}) \wedge \\
& \text{are\_concrete\_flyweight}(\text{dr}) \wedge \\
& \text{FF\_stores\_flyweights}(\text{dr}) \wedge \\
& \text{FF\_aggregates\_flyweights}(\text{dr}) \wedge \\
& \text{FF\_creates\_flyweights}(\text{dr}) \wedge \\
& \text{F\_has\_defined\_operation}(\text{dr}) \wedge \\
& \text{CF\_has\_impl\_operation}(\text{dr}) \wedge \\
& \text{UCF\_has\_impl\_operation}(\text{dr}) \wedge \\
& \text{CF\_stores\_intrinsic\_state}(\text{dr}) \wedge \\
& \text{UCF\_stores\_all\_state}(\text{dr}) \wedge \\
& \text{FF\_has\_impl\_get\_flyweight}(\text{dr}) \wedge \\
& \text{FF\_has\_impl\_create\_flyweight}(\text{dr})
\end{aligned}$$

## 8 The Proxy Pattern

The Proxy pattern offers a way of providing a reference (proxy) to an object which is more versatile and more sophisticated than a simple pointer. A proxy can be used, for example, to provide a local representative for an object in a different address space (*remote proxy*), to defer creation of large objects (for instance graphical objects) until they are needed (*virtual proxy*), to control access to an object (*protection proxy*), or to perform additional actions when an object is accessed (*smart reference*).

Proxies offer interfaces which are identical to those of the objects they represent, so clients never notice that an object has been substituted by a proxy. Proxies thus provide a kind of indirect pointer to an object.

We begin as usual by introducing the essential elements of the Proxy pattern as defined in [7], then we give a formal specification of these properties. Finally, we describe and specify a version of the pattern based on one of the specific kinds of proxy mentioned above, the virtual proxy.

### 8.1 Properties of the Proxy Pattern

#### Intent

Provide a surrogate or placeholder for another object to control access to it.

#### Structure

Figure 16 shows the structure of the Proxy pattern, and Figure 17 shows an object diagram of a typical proxy structure in run-time.

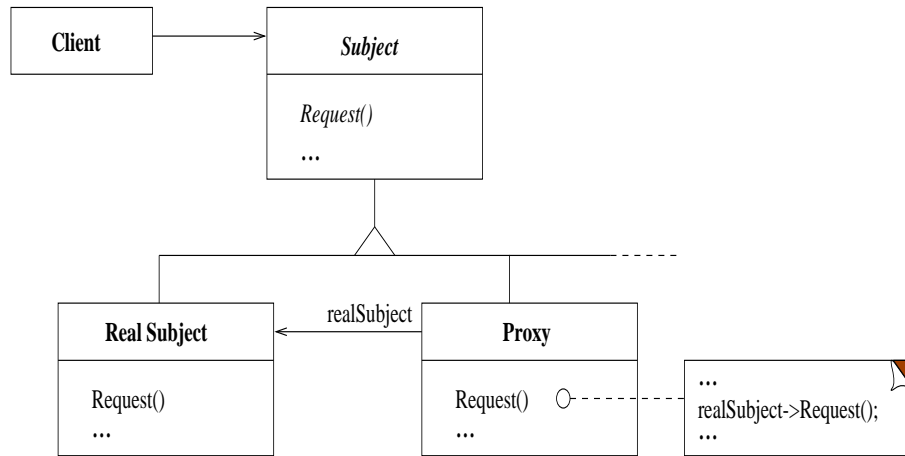


Figure 16: Proxy Pattern Structure




Figure 17: Proxy Pattern Object Structure

## Participants

- Proxy
  - maintains a reference that lets the proxy access the real subject. Proxy may refer to a Subject if the RealSubject and Subject interfaces are the same.
  - provides an interface identical to Subject's so that a proxy can be substituted for the real subject.
  - controls access to the real subject and may be responsible for creating and deleting it.
  - other responsibilities depend on the kind of proxy:
    - \* *remote proxies* are responsible for encoding a request and its arguments and for sending the encoded request to the real subject in a different address space.
    - \* *virtual proxies* may cache additional information about the real subject so that they can postpone accessing it.
    - \* *protection proxies* check that the caller has the access permissions required to perform a request.
- Subject

- defines the common interface for RealSubject and Proxy so that a Proxy can be used anywhere a RealSubject is expected.
- RealSubject
  - defines the real object that the proxy represents.

### Collaborations

- Proxy forwards requests to RealSubject when appropriate, depending on the kind of proxy.

## 8.2 Formalising the Proxy Pattern

Before beginning the formal specification of the Proxy pattern, we note that the pattern structure shown in Figure 16 is in fact incomplete in general – there can be more than two subclasses of the Subject class as indicated by the dashed line in the right-hand part of the inheritance relationship. We therefore make this possibility explicit by introducing a new class, which we call ConcreteSubject, to the pattern. Figure 18 shows the modified pattern structure.

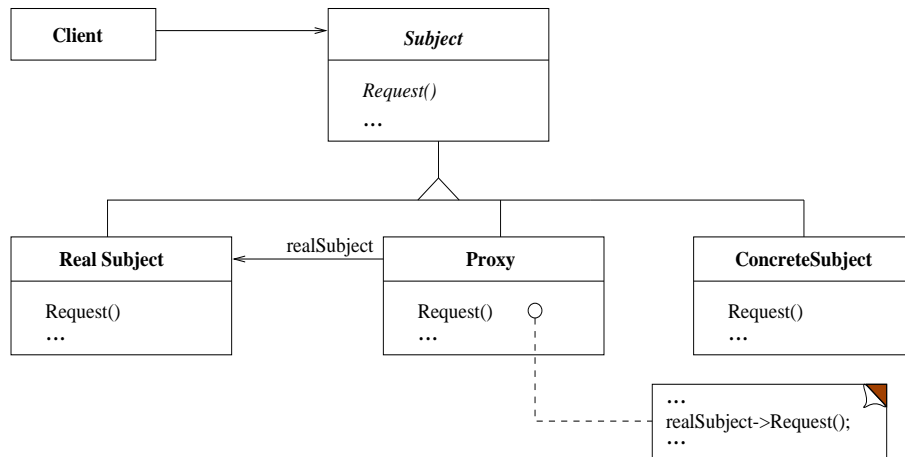


Figure 18: The Modified Proxy Pattern Structure

The entities in this modified pattern structure are represented by the following RSL constants:

#### value

Subject, RealSubject, Proxy, ConcreteSubject, Client : G.Class\_Name,

realSubject : G.Variable\_Name,

Request : G.Method\_Name

The structure essentially consists of a single hierarchy of classes with the **Subject** role at the root of the hierarchy and the roles **RealSubject**, **Proxy** and **ConcreteSubject** at the leaves. Again we can consider this hierarchy, and hence the **Subject** role, to be unique in the design without any loss of generality, assuming as usual that each such hierarchy in the design corresponds to a distinct instance of the pattern. We also require that the three roles **RealSubject**, **Proxy** and **ConcreteSubject** are distinct, that is that no class in the design plays more than one of these roles, since the roles have different responsibilities. Furthermore, no class in the hierarchy can play the **Client** role.

The properties of this hierarchy are therefore specified in the standard way using the function *hierarchy*.

**value**

```

one_subject_in_hierarchy : Wf_Design_Renaming → Bool
one_subject_in_hierarchy(dr) ≡
  hierarchy
    (Subject, {RealSubject, Proxy, ConcreteSubject}, {Client}, dr)

```

Each object of the **Proxy** class represents an object of the **RealSubject** class in such a way that the difference is transparent to clients, i.e. clients cannot tell whether they are interacting with the **Proxy** or the **RealSubject** class. This is achieved by ensuring that the interfaces of the two classes **Proxy** and **RealSubject** are the same, as indicated in the structure. However, the structure only shows one class playing the **RealSubject** role whereas in practice this class could be a superclass which defines the common interface for a whole hierarchy of **RealSubject** classes as in the example shown in part A of Figure 19, and in this case an object of the **Proxy** class could equally represent an object from any of the **RealSubject** classes in the hierarchy.

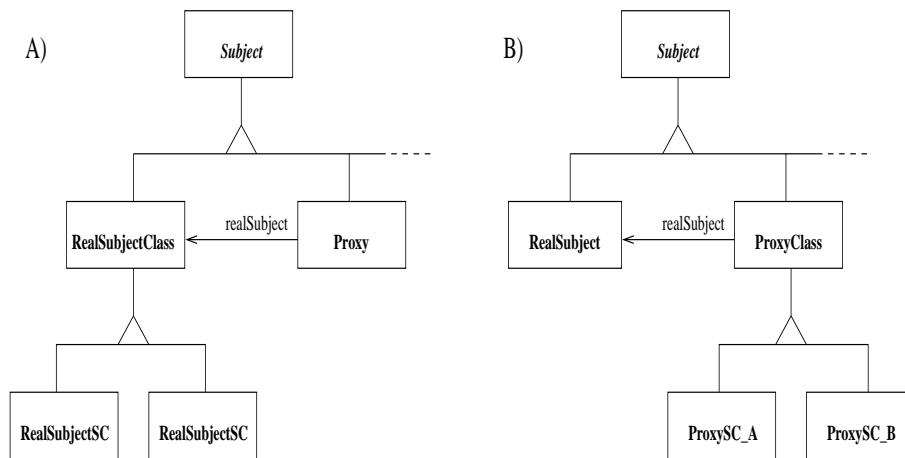


Figure 19: RealSubject and Proxy Hierarchies in the Proxy Pattern

Another possibility might be that we have a similar hierarchy of **Proxy** classes instead of a single such class, as shown in part B of Figure 19. However, in this case it does not make sense to have

a single relation linking the root of the **Proxy** hierarchy to the root of the **RealSubject** hierarchy as shown in the figure since this would correspond to a situation in which an object from any of the **Proxy** classes could represent an object from any of the **RealSubject** classes. This effectively gives clients many different possible surrogates for any single type of object, which is not the purpose of the pattern. We therefore require that there is no hierarchy of **Proxy** classes.

In fact we can further restrict to a single class in the design which plays the **Proxy** role and a single hierarchy of classes playing the **RealSubject** role, as usual considering without any loss of generality that each such combination corresponds to a separate instance of the pattern. We specify the uniqueness of the **Proxy** class in the standard way using the function *exists\_one*, and the existence and uniqueness of the single hierarchy of **RealSubject** classes using the functions *exists\_role* and *only\_one\_hierarchy* from [6].

```

value
  exists_real_subject : Wf_Design_Renaming → Bool
  exists_real_subject(dr) ≡
    exists_role(RealSubject, dr) ∧
    only_one_hierarchy(RealSubject, dr),

  exists_one_proxy : Wf_Design_Renaming → Bool
  exists_one_proxy(dr) ≡ exists_one(Proxy, dr)

```

Of course, the class playing the **Proxy** role and all classes playing the **RealSubject** role must be concrete subclasses of the class playing the **Subject** role. In addition, any class playing the **ConcreteSubject** role must also be a concrete subclass of the class playing the **Subject** role, though in this case it is not necessary that the design actually includes such classes. We specify all these properties using the function *is\_concrete* as usual.

```

value
  are_concrete_subject : Wf_Design_Renaming → Bool
  are_concrete_subject(dr) ≡
    is_concrete(Subject, RealSubject, dr) ∧
    is_concrete(Subject, Proxy, dr) ∧
    is_concrete(Subject, ConcreteSubject, dr)

```

In order to carry out its responsibilities as surrogate, a proxy stores a reference to the object it represents in its *realSubject* state variable. Only one such variable is needed since a proxy represents a single object. There is also a one-one association relation corresponding to this state variable which links the **Proxy** class to the **RealSubject** class. This relation is also unique (that is no more association or aggregation relations between these two classes are required). However, when the design contains a hierarchy of **RealSubject** classes, the association relation must link to the root of this hierarchy instead of to some class in the middle so that the single

Proxy class can reasonably represent any of the `RealSubject` classes as explained above. This last property is embodied in the function *has\_ass\_agg\_var\_ren\_one\_sink* from [6], while the existence and uniqueness of the state variable and the relation are specified in the usual way using the functions *store\_unique\_vble* and *has\_unique\_assoc\_aggr\_relation* respectively.

Note that the description of the participants of the pattern given in Section 8.1 states that the Proxy class may refer to the Subject class if the `RealSubject` and `Subject` interfaces are the same. Indeed, [14] presents an example using a combination of the Proxy pattern and the Composite pattern in which the Proxy class refers directly to the Subject class. However, although this situation is certainly possible, it has a structure which is different from that shown in Figure 18 so we consider it to represent a variant of the pattern. We intend to consider such variants in future work.

**value**

```
PX_stores_one_real_subject : Wf_Design_Renaming → Bool
PX_stores_one_real_subject(dr) ≡ store_unique_vble(Proxy, realSubject, dr),

PX_references_RS : Wf_Design_Renaming → Bool
PX_references_RS(dr) ≡
  has_ass_agg_var_ren_one_sink
    (Proxy, RealSubject, Association, realSubject, G.one, dr) ∧
  has_unique_assoc_aggr_relation(Proxy, RealSubject, dr)
```

The `Subject` class defines the interface used by clients, and this is represented by the `Request` method. In general this interface can of course consist of many methods, each with a different functionality, so there are no restrictions on the methods except that they must exist.

At the level of the `Subject` class, the interface is only defined (abstract), so we once again use the functions *has\_def\_method* and *has\_all\_def\_method* to specify it as in the specification of the properties of the `Request` method in the `Target` class in both versions of the Adapter pattern.

**value**

```
S_has_defined_request : Wf_Design_Renaming → Bool
S_has_defined_request(dr) ≡
  has_def_method(Subject, Request, dr) ∧
  has_all_def_method(Subject, Request, dr)
```

Since the client uses the `Subject` class interface to interact with the Proxy class and through that the `RealSubject` class, these two classes must implement the `Request` interface. In the `RealSubject` class the implementation is unconstrained, but the annotation to the `Request` method in the Proxy class indicates that it includes an invocation to the `realSubject` state variable of the same `Request` method – this is the means by which the proxy acts as surrogate for the real subject.

These properties are again specified using the functions *has\_all\_impl\_method* and *deleg\_with\_var*. For the **ConcreteSubject** classes, there are no restrictions on the **Request** methods – some may be implemented but some may be error methods. This property follows automatically from the fact that the **ConcreteSubject** classes are concrete subclasses of the **Subject** class so we do not need to specify anything else here.

**value**

```

RS_has_impl_request : Wf_Design_Renaming → Bool
RS_has_impl_request(dr) ≡ has_all_impl_method(RealSubject, Request, dr),

PX_has_impl_request : Wf_Design_Renaming → Bool
PX_has_impl_request(dr) ≡
  deleg_with_var(Proxy, Request, realSubject, RealSubject, Request, dr)

```

In order to make the difference between the **RealSubject** and the **Proxy** transparent to clients, the **Client** class interacts solely with the abstract **Subject** class which provides a common interface to those two classes. This interaction is through the **Request** methods in the interface of the **Subject** class and it is represented in the structure by the relation linking the **Client** to the **Subject** class, which may be either an association or an aggregation relation depending on circumstances but which in any case has cardinality one-one.

These properties are exactly analogous to the properties of the **Client** role in the Bridge pattern (see Section 3.2) so they are also formally specified using the functions *has\_assoc\_aggr\_reltype* and *use\_interface*.

**value**

```

proxy_client : Wf_Design_Renaming → Bool
proxy_client(dr) ≡
  has_assoc_aggr_reltype(Client, Subject, AssAggr, G.one, dr) ∧
  use_interface(Client, Subject, Request, dr)

```

The function *is\_proxy\_pattern* combines all the above properties and thus can be used to check whether or not a given design matches the Proxy pattern.

**value**

```

is_proxy_pattern : Wf_Design_Renaming → Bool
is_proxy_pattern(dr) ≡
  one_subject_in_hierarchy(dr) ∧
  exists_real_subject(dr) ∧
  exists_one_proxy(dr) ∧
  proxy_client(dr) ∧

```

```

are_concrete_subject(dr) ∧
PX_has_identical_S_intf(dr) ∧
PX_stores_one_real_subject(dr) ∧
PX_references_RS(dr) ∧
S_has_defined_request(dr) ∧
PX_has_impl_request(dr) ∧
RS_has_impl_request(dr)

```

### 8.3 Virtual Proxy

The analysis and specification presented above relate to the general Proxy pattern, but as explained earlier there are several different kinds of proxies, each of which has some particular additional properties and behaviour as outlined in the description of the pattern's participants. In this section we illustrate how to extend our general specification to obtain a specification of one of these kinds of proxy, the virtual proxy.

A virtual proxy allows the management of objects to take place at the appropriate time, i.e. when they are actually needed. For example, clients sometimes need to use objects that have large, complex objects associated with them, and if these associated objects are not necessarily required it is better to delay their creation until they are actually needed to avoid the overhead of creating and storing them in situations where they are unnecessary. In such situations, the large, complex object is created on demand and another object, the virtual proxy, acts as a surrogate for the object, controlling access to it and taking care of instantiating it when it is actually required.

In the case of the virtual proxy, therefore, the Proxy class must be able to create the RealSubject object as well as forward requests to it from the Client. We make this new behaviour explicit by introducing a new method role, **Creator**, into the Proxy class, and we extend the formal specification by adding a new RSL constant to represent this new role:

**value**

Creator : G.Method\_Name

In order to create the RealSubject object, this method should be implemented and its body should include an instantiation of a RealSubject class, the result of this instantiation being the result of the method. This second property is specified by the function *has\_method\_with\_result\_class*, while the function *has\_impl\_method* is used as usual to ensure that the **Creator** method is concrete.

**value**

PX\_has\_impl\_creator : Wf\_Design\_Renaming → **Bool**



```

PX_has_impl_creator(dr)  $\equiv$ 
  has_impl_method(Proxy, Creator, dr)  $\wedge$ 
  has_method_with_result_class(Proxy, Creator, RealSubject, dr)

```

Since the `Creator` method creates objects of the `RealSubject` class, there should be an instantiation relation linking the `Proxy` class to the `RealSubject` classes. Again in the case in which there is a hierarchy of `RealSubject` classes as shown in part A of Figure 19, this relation should link to the root of the hierarchy. We specify the existence of this relation using the function *has\_instantiation*.

**value**

```

PX_creates_RS : Wf_Design_Renaming  $\rightarrow$  Bool
PX_creates_RS(dr)  $\equiv$  has_instantiation(RealSubject, Proxy, dr)

```

Although there could be a single method in the design playing the `Creator` role, this method being able to create objects belonging to every `RealSubject` class, some designs may have specific methods for different kinds of real subject. In any case, the `Creator` methods should not be accessible directly by the `Client` and in fact should only be invoked from within the `Proxy` class. The `Creator` methods are thus part of the private interface of the `Proxy` class. This property is expressed formally using the function *has\_private\_interface* from [6].

**value**

```

PX_has_private_creator : Wf_Design_Renaming  $\rightarrow$  Bool
PX_has_private_creator(dr)  $\equiv$  has_private_interface(Proxy, Creator, dr)

```

Then, since a virtual proxy is a basic `Proxy` pattern with the additional properties described and specified above, a design corresponds to an instance of the “Virtual Proxy” pattern if it satisfies the function *is\_proxy\_pattern* together with the functions specifying these additional properties. The function *is\_virtual\_proxy\_pattern* captures this.

**value**

```

is_virtual_proxy_pattern : Wf_Design_Renaming  $\rightarrow$  Bool
is_virtual_proxy_pattern(dr)  $\equiv$ 
  is_proxy_pattern(dr)  $\wedge$ 
  PX_creates_RS(dr)  $\wedge$ 
  PX_has_impl_creator(dr)  $\wedge$ 
  PX_has_private_creator(dr)

```

## 9 Application of the Decorator Pattern

In order to show how the formal specifications presented in this report can be used to check whether a given design corresponds to an instance of some structural pattern, we consider in this section a design for a “decorated text editor” which in fact corresponds to an application of the Decorator pattern. This example is in fact based on the example discussed in [7] as part of the motivation of the Decorator pattern.

The example is concerned with the design of a graphical user interface for a text editor in which it is necessary to add in a flexible way two embellishments for displaying the text. The first adds a border around the text editing area to demarcate the page of text, while the second adds scroll bars that let the user view different parts of the page. The extended OMT diagram for this design is shown in Figure 20.

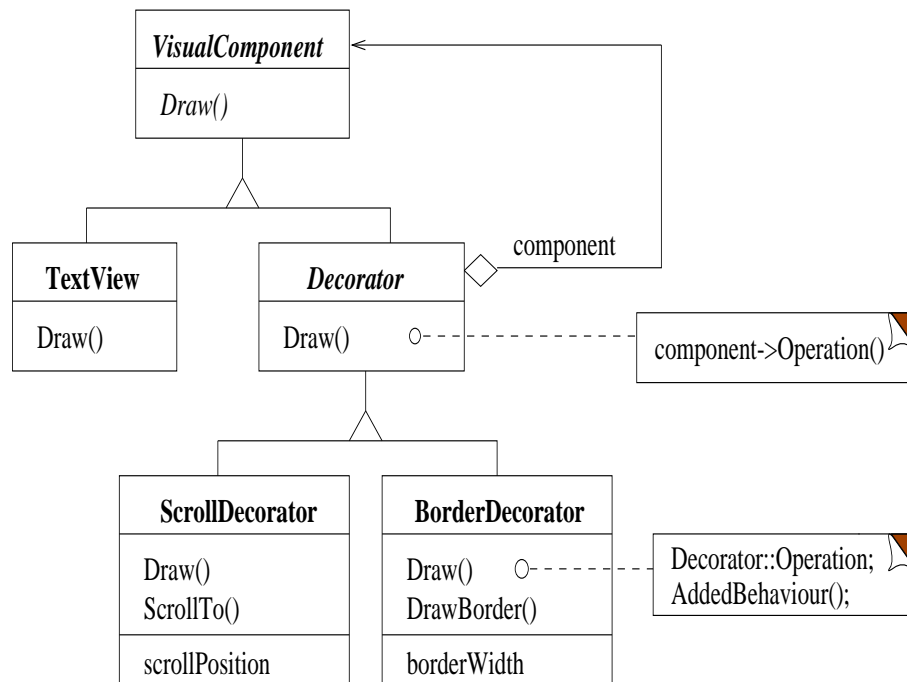


Figure 20: Design for a Decorated Text Editor

The class `TextView` represents the objects that display text in a window. By default, the `TextView` class does not include scroll bars or borders since these do not necessarily appear on all views. When they are required, the `ScrollDecorator` or `BorderDecorator` class is used to add them respectively, and any desired combination can be obtained by composing a `TextView` with a `ScrollDecorator` and/or a `BorderDecorator` as appropriate. For example, an object composition producing a bordered, scrollable text view is shown in the object diagram in Figure 21.

In order to show that the design in Figure 20 is actually an instance of the Decorator pattern, we first specify the properties of all the entities in the design using the general model defined in [6],




Figure 21: Decorated Text View Object Diagram

and we check that the design is well-formed by checking that the well-formedness conditions on the model are satisfied. Next we specify the renaming map which relates the names of the entities appearing in the design to the corresponding names in the appropriate pattern (Decorator). Finally, we check that the specification of the design satisfies the defining function for the pattern, which in the case of the Decorator pattern is the function *is\_decorator\_pattern*. We only give an overview of the specification here. The full details can be found in Appendix A.

The design includes five classes, three methods, and three variables, and we introduce a sixth class, *User*, to represent the users or clients of the text editor system. Each of these entities is then represented by an RSL constant of the appropriate type:

**value**

VisualComponent, TextView, Decorator, ScrollDecorator,  
BorderDecorator, User : G.Class\_Name,

component, scrollPosition, borderWidth : G.Variable\_Name,

Draw, ScrollTo, DrawBorder : G.Method\_Name

Each class is described by specifying an appropriate instance of the type *Design\_Class*, which is a record type comprising the state variables and operations belonging to the class together with its type (abstract or concrete). For example, the class *VisualComponent* has no state variables and a single (abstract) method called *Draw*, and its type is *abstract*.

Every operation is similarly described by specifying an appropriate instance of the record type *Class\_Method*, which comprises the formal parameters of the method, its result, and its body. The method *Draw* has no formal parameters, no result and it is abstract so its body is simply the constant *defined*.

The RSL constant definitions *VC\_Design\_Class* and *VC\_Class\_Method* represent the class *VisualComponent* and the method *Draw* respectively.

**value**

VC\_Class\_Method : M.Class\_Method =  
[Draw  $\mapsto$  M.mk\_Method( $\langle \rangle$ ,  $\{ \}$ , M.defined)],

```

VC_Design_Class : C.Design_Class =
  C.mk_Design_Class({}, VC_Class_Method, G.abstract)

```

An implemented operation may include assignments to local variables or to state variables in its *variable\_change* mapping, these assignments involving either other variables or parameters or the results from internal invocations of other methods. In addition there may be a list of interactions the method performs, which is modelled as its *request\_list*, each element of which can be an invocation, an instantiation, etc.

The `Decorator` class implements the `Draw` operation, and the body of this operation contains an invocation to the `component` state variable of the `Draw` method in the `VisualComponent` class. The RSL constant *D\_Design\_Class* defines the properties of the `Decorator` class, namely that it is abstract and includes a state variable called `component` and the implemented operation `Draw`. The properties of the method `Draw` are similarly embodied in the constant *D\_Class\_Method*, with the constant *meth\_body\_Dw\_D* defining the properties of the body of the method, that is that the method has an empty variable change map and a single invocation in its request list which forwards the `Draw` method to the `component` state variable.

**value**

```

D_Design_Class : C.Design_Class =
  C.mk_Design_Class({component}, D_Class_Method, G.abstract),

D_Class_Method : M.Class_Method =
  [Draw ↦ M.mk_Method(⟨⟩, {}, meth_body_Dw_D)],

meth_body_Dw_D : M.Method_Body =
  M.implemented([],
    ⟨M.mk_Invocation(component, M.mk_Actual_Signature(Draw,⟨⟩))⟩
  )

```

The body of the `Draw` method in the `BorderDecorator` class includes a super invocation to the `Draw` method in the `Decorator` class followed by a self invocation of the `DrawBorder` method. These properties are captured in the constant *meth\_body\_Dw\_BD*, which specifies that the method has no *variable\_change* but its body includes the two appropriate invocations:

**value**

```

meth_body_Dw_BD : M.Method_Body =
  M.implemented([],
    ⟨M.mk_Invocation(G.super, M.mk_Actual_Signature(Draw,⟨⟩)),
      M.mk_Invocation(G.self, M.mk_Actual_Signature(DrawBorder,⟨⟩))⟩
  )

```

The rest of the classes in the design can be specified in a similar way. Full details can be found in Appendix A. Then the collection of all classes in the design is specified as a mapping from the name of each class to its definition as follows:

```

value
Deco_Classes : C.Classes =
  [ VisualComponent  $\mapsto$  VC_Design_Class,
    TextView  $\mapsto$  TV_Design_Class,
    Decorator  $\mapsto$  D_Design_Class,
    ScrollDecorator  $\mapsto$  SD_Design_Class,
    BorderDecorator  $\mapsto$  BD_Design_Class ]

```

Relationships are similarly specified as instances of the type *Design\_Relation*. This is again a record type comprising the names of the classes related (source and sink) together with the type of the relation (inheritance, instantiation, aggregation or association). For aggregation and association relations the name of the corresponding variable and the cardinality of the relation (one or many) are also given.

Thus, for example, the inheritance relation between the *VisualComponent* and the *TextView* classes is defined by the constant *VC\_inh\_TV* as follows:

```

value
VC_inh_TV : R.Design_Relation =
  R.mk_Design_Relation(R.inheritance, VisualComponent, TextView)

```

Similarly, the aggregation relation between the *Decorator* and *VisualComponent* classes is formally expressed as the constant *D\_aggr\_VC*, the type, name and cardinality of the relation being represented by the constant *D\_rel\_type*:

```

value
D_aggr_VC : R.Design_Relation =
  R.mk_Design_Relation(D_rel_type, Decorator, VisualComponent),

D_rel_type : R.Relation_Type =
  R.aggregation(R.mk_Ref(component, G.one, G.one))

```

Then a new constant *Deco\_Relations* is introduced to represent the set of all relations in the design, and the whole structure of the design is represented by the instance of the type *Design\_Structure* obtained by combining the classes and relations (after checking that the appropriate well-formedness conditions on the classes and relations and all their constituent parts are satisfied).

**value**

```
Deco_Relations : R.Wf_Relation-set =
  {VC_inh_TV, VC_inh_D, D_inh_SD, D_inh_BD, D_aggr_VC},

Deco_Structure : DS.Design_Structure = (Deco_Classes, Deco_Relations)
```

The next step is to define the link between the entities in the design and the entities in the Decorator pattern by constructing the renaming map, which is an instance of the type *Renaming*. That is, we must assign appropriate pattern roles to the entities in the design.

We associate the method `Draw` with the `Operation` role via the method renaming *VC\_metd\_ren*. This has an empty parameter renaming since the method has no formal parameters. Similarly, we associate the `VisualComponent` class with the `Component` role through the class renaming *VC\_ClassRenaming*. Here, the renaming for state variables is empty because the `VisualComponent` class has an empty state and the method renaming *VC\_metd\_ren* describes the renaming of the method `Draw` within the class.

**value**

```
VC_metd_ren : Method_and_Parameter_Renaming =
  [ Draw  $\mapsto$  mk_Method_Renaming(DEC.Operation, [])],

VC_ClassRenaming : ClassRenaming =
  mk_ClassRenaming(DEC.Component, VC_metd_ren, [])
```

The component state variable in the Decorator class plays the component role in the pattern, and the method renaming *VC\_metd\_ren* again describes the renaming of the method `Draw` within the class. The renaming of the Decorator class is thus defined by the class renaming *D\_ClassRenaming*.

**value**

```
D_var_ren : VariableRenaming = [component  $\mapsto$  DEC.component],

D_ClassRenaming : ClassRenaming =
  mk_ClassRenaming(DEC.Decorator, VC_metd_ren, D_var_ren)
```

The `ScrollDecorator` class plays the `ConcreteDecorator` role, with the `scrollPosition` variable playing the `addedState` role and the `ScrollTo` method playing the `AddedBehaviour` role. The following declarations specify this part of the renaming:

**value**

```

SD_var_ren : VariableRenaming = [scrollPosition ↦ DEC.addedState],

SD_metd_ren : Method_and_Parameter_Renaming =
  [ Draw ↦ mk_Method_Renaming(DEC.Operation, []),
    ScrollTo ↦ mk_Method_Renaming(DEC.AddedBehaviour, []) ],

SD_ClassRenaming : ClassRenaming =
  mk_ClassRenaming(DEC.ConcreteDecorator, SD_metd_ren, SD_var_ren)

```

The renamings of the other design classes can be specified in a similar way. Again, full details can be found in Appendix A.

The renaming for all the classes in the design is then formed by simply combining the individual class renamings into the map *Deco\_Renaming*, and further combining this with the constant *Deco\_Structure* which represents the design structure yields the constant *Deco\_Design\_Renaming*. (Again we must check first that all appropriate well-formedness conditions on the renaming and the combination of the renaming and the structure are satisfied.)

```

value
Deco_Renaming : Renaming =
  [ VisualComponent ↦ {VC_ClassRenaming},
    TextView ↦ {TV_ClassRenaming},
    Decorator ↦ {D_ClassRenaming},
    ScrollDecorator ↦ {SD_ClassRenaming},
    BorderDecorator ↦ {BD_ClassRenaming},
    User ↦ {Usr_ClassRenaming} ],

Deco_Design_Renaming : Wf_Design_Renaming =
  (Deco_Structure, Deco_Renaming)

```

This value is then used as input to the function *is\_decorator\_pattern* defined in Section 5.2 to check whether or not the decorated text editor design is an instance of the Decorator pattern.

## 10 Conclusions

In this report we have presented an analysis and formal specification of the essential properties of each of the GoF structural patterns using the model of an object-oriented design developed in earlier work [6]. We have also illustrated how the process of binding a design to a pattern is carried out using an example based on the Decorator pattern, and we have further indicated how the model can be used to check that a design matches a particular pattern.

Our analysis has also led to the identification of several ambiguities and incompletenesses in the textual/graphical descriptions of a number of the structural patterns, as a result of which we have proposed modifications to the structures of these patterns which resolve these problems.

We believe this formal approach to GoF structural patterns can be a useful complement to the commonly used informal notation, helping designers to improve their understanding of the patterns and also giving a means of checking formally that designs match patterns or that patterns are being applied correctly. We also believe that the model could form the basis for software tools supporting the use of patterns and we plan to investigate this in future work.

## A Specification of the Decorated Text Editor

In the following specification, items prefixed with the object name ‘DEC’ refer to declarations from the specification of the Decorator pattern given in Section 5.

**scheme**

DECORATOR\_EXA =

**class**

**value**

VisualComponent, TextView, Decorator, ScrollDecorator,  
BorderDecorator, User : G.Class\_Name,

component, scrollPosition, borderWidth : G.Variable\_Name,

Draw, ScrollTo, DrawBorder : G.Method\_Name,

VC\_metd\_ren : Method\_and\_Parameter\_Renaming =  
[ Draw  $\mapsto$  mk\_Method\_Renaming(DEC.Operation, [])],

VC\_ClassRenaming : ClassRenaming =  
mk\_ClassRenaming(DEC.Component, VC\_metd\_ren, []),

TV\_ClassRenaming : ClassRenaming =  
mk\_ClassRenaming(DEC.ConcreteComponent, VC\_metd\_ren, []),

D\_var\_ren : VariableRenaming = [component  $\mapsto$  DEC.component],

D\_ClassRenaming : ClassRenaming =  
mk\_ClassRenaming(DEC.Decorator, VC\_metd\_ren, D\_var\_ren),

SD\_var\_ren : VariableRenaming = [scrollPosition  $\mapsto$  DEC.addedState],



```

SD_metd_ren : Method_and_Parameter_Renaming =
  [ Draw ↦ mk_Method_Renaming(DEC.Operation, []),
    ScrollTo ↦ mk_Method_Renaming(DEC.AddedBehaviour, []) ],

SD_ClassRenaming : ClassRenaming =
  mk_ClassRenaming(DEC.ConcreteDecorator, SD_metd_ren, SD_var_ren),

BD_var_ren : VariableRenaming = [ borderWidth ↦ DEC.addedState ],

BD_metd_ren : Method_and_Parameter_Renaming =
  [ Draw ↦ mk_Method_Renaming(DEC.Operation, []),
    DrawBorder ↦ mk_Method_Renaming(DEC.AddedBehaviour, []) ],

BD_ClassRenaming : ClassRenaming =
  mk_ClassRenaming(DEC.ConcreteDecorator, BD_metd_ren, BD_var_ren),

Usr_ClassRenaming : ClassRenaming =
  mk_ClassRenaming(DEC.Client, [], []),

Deco_Renaming : Renaming =
  [ VisualComponent ↦ {VC_ClassRenaming},
    TextView ↦ {TV_ClassRenaming},
    Decorator ↦ {D_ClassRenaming},
    ScrollDecorator ↦ {SD_ClassRenaming},
    BorderDecorator ↦ {BD_ClassRenaming},
    User ↦ {Usr_ClassRenaming} ],

/* Decorator Draw method Body */
meth_body_Dw_D : M.Method_Body =
  M.implemented([],
    ⟨M.mk_Invocation(component,
      M.mk_Actual_Signature(Draw, ⟨⟩))),

meth_body_Dw_SD : M.Method_Body =
  M.implemented([],
    ⟨M.mk_Invocation(G.super,
      M.mk_Actual_Signature(Draw, ⟨⟩),
      M.mk_Invocation(G.self,
        M.mk_Actual_Signature(ScrollTo, ⟨⟩))),

meth_body_Dw_BD : M.Method_Body =
  M.implemented([],
    ⟨M.mk_Invocation(G.super,
      M.mk_Actual_Signature(Draw, ⟨⟩),

```

```

M.mk_Invocation(G.self,
  M.mk_Actual_Signature(DrawBorder, ⟨⟩)),

meth_body_impl : M.Method_Body = M.implemented([], ⟨⟩),

D_Class_Method : M.Class_Method =
  [ Draw ↦ M.mk_Method(⟨⟩, {}, meth_body_Dw_D) ],

SD_Class_Method : M.Class_Method =
  [ Draw ↦ M.mk_Method(⟨⟩, {}, meth_body_Dw_SD),
    ScrollTo ↦ M.mk_Method(⟨⟩, {}, meth_body_impl) ],

BD_Class_Method : M.Class_Method =
  [ Draw ↦ M.mk_Method(⟨⟩, {}, meth_body_Dw_BD),
    DrawBorder ↦ M.mk_Method(⟨⟩, {}, meth_body_impl) ],

VC_Class_Method : M.Class_Method =
  [ Draw ↦ M.mk_Method(⟨⟩, {}, M.defined) ],

TV_Class_Method : M.Class_Method =
  [ Draw ↦ M.mk_Method(⟨⟩, {}, meth_body_impl) ],

SD_Design_Class : C.Design_Class =
  C.mk_Design_Class({scrollPosition}, SD_Class_Method, G.concrete),

BD_Design_Class : C.Design_Class =
  C.mk_Design_Class({borderWidth}, SD_Class_Method, G.concrete),

D_Design_Class : C.Design_Class =
  C.mk_Design_Class({component}, D_Class_Method, G.abstract),

TV_Design_Class : C.Design_Class =
  C.mk_Design_Class({}, TV_Class_Method, G.concrete),

VC_Design_Class : C.Design_Class =
  C.mk_Design_Class({}, VC_Class_Method, G.abstract),

Deco_Classes : C.Classes =
  [ VisualComponent ↦ VC_Design_Class,
    TextView ↦ TV_Design_Class,
    Decorator ↦ D_Design_Class,
    ScrollDecorator ↦ SD_Design_Class,

```

```

    BorderDecorator  $\mapsto$  BD_Design_Class],

VC_inh_TV : R.Design_Relation =
    R.mk_Design_Relation(R.inheritance, VisualComponent, TextView),

VC_inh_D : R.Design_Relation =
    R.mk_Design_Relation(R.inheritance, VisualComponent, Decorator),

D_inh_SD : R.Design_Relation =
    R.mk_Design_Relation(R.inheritance, Decorator, ScrollDecorator),

D_inh_BD : R.Design_Relation =
    R.mk_Design_Relation(R.inheritance, Decorator, BorderDecorator),

D_rel_type : R.Relation_Type =
    R.aggregation(R.mk_Ref(component, G.one, G.one)),

D_aggr_VC : R.Design_Relation =
    R.mk_Design_Relation(D_rel_type, Decorator, VisualComponent),

Deco_Relations : R.Wf_Relation-set =
    {VC_inh_TV, VC_inh_D, D_inh_SD, D_inh_BD, D_aggr_VC},

Deco_Structure : DS.Design_Structure =
    (Deco_Classes, Deco_Relations),

Deco_Design_Renaming : Wf_Design_Renaming = (Deco_Structure, Deco_Renaming)

end

```

## References

- [1] Brad Appleton. Patterns and Software: Essential Concepts and Terminology. <http://www.enteract.com/~bradapp>, November 1997.
- [2] Gabriela Aranda and Richard Moore. GoF Creational Patterns: A Formal Specification. Technical Report 224, UNU/IIST, P.O. Box 3058, Macau, December 2000.

- [3] N. Martinez Carod, A. Flores, and L. Reynoso. A Formal Specification of a Pattern based on Recursive Structure: The Composite Pattern. Technical report nro. 001, UNC/AIS, Neuquen - Argentina, 1999.
- [4] S. Alejandra Cechich and Richard Moore. A Formal Specification of GoF Design Patterns. In *Proceedings of the Asia Pacific Software Engineering Conference: APSEC'99*, Takamatsu, Japan, December 1999.
- [5] Amnon H. Eden. Giving The Quality a Name. *Journal of Object Oriented Programming*, May 1997. <http://www.math.tau.ac.il/~eden/bibliography>.
- [6] Andres Flores, Luis Reynoso, and Richard Moore. A Formal Model of Object-Oriented Design and GoF Design Patterns. Technical Report 200, UNU/IIST, P.O. Box 3058, Macau, July 2000. To be presented at and published in the proceedings of FME 2001, Berlin, Germany, 12-16 March 2001.
- [7] Erich Gamma, Richard Helm, Ralph Johnson, and John Vlissides. *Design Patterns: Elements of Reusable Object-Oriented Software*. Addison Wesley Professional Computing Series. Addison Wesley, 1995.
- [8] Doug Lea. Patterns-Discussion FAQ. <http://g.oswego.edu/dl/pd-FAQ/pd-FAQ.html>, December 1999.
- [9] James Martin and James J. Odell. *Object-Oriented Methods: A Foundation*. Prentice Hall, 1995.
- [10] W. Pree. *Design Patterns for Object-Oriented Software Development*. Addison-Wesley, 1995.
- [11] The RAISE Language Group. *The RAISE Specification Language*. BCS Practitioner Series. Prentice Hall, 1992.
- [12] Luis Reynoso and Richard Moore. GoF Behavioural Patterns: A Formal Specification. Technical Report 201, UNU/IIST, P.O. Box 3058, Macau, May 2000.
- [13] J. Rumbaugh, M. Blaha, W. Premerlani, F. Eddy, and W. Lorensen. *Object-Oriented Modelling and Design*. Prentice-Hall, 1991.
- [14] John Vlissides. Notation, Notation, Notation. C++ Report, April 1998. <http://www.research.ibm.com/people/v/vlis/pubs.html>.
- [15] B. Wydaeghe, K. Verschaeve, B. Michiels, B. Van Damme, E. Arckens, and V. Jonckers. Building an OMT-Editor Using Design Patterns: An Experience Report. Vrije Universiteit Brussel, 1998. <http://info.vub.ac.be/~bwydaegh/tekst/paers/tools98/html/tools98.html>.
- [16] Robert Zubeck. Much Ado About Patterns . <http://www.acm.org/crossroads/xrds5-1/patterns.html>, March 2000.